

# No-Regret Learning for Nonsubmodular Functions: Bridging Full Information and Bandit Feedback

Chris Junchi Li<sup>◊</sup>

Department of Electrical Engineering and Computer Sciences<sup>◊</sup>  
University of California, Berkeley

September 2, 2024

## Abstract

In this paper, we present novel adaptive algorithms for minimizing nonsubmodular functions in online learning environments, where decision-making is challenged by delayed feedback and incomplete information. Our approach extends traditional online learning frameworks by incorporating a randomized gradient descent method tailored for nonsubmodular cost functions. We introduce a new notion of  $(\alpha, \beta)$ -regret to evaluate the performance of these algorithms in both full-information and bandit settings. Our theoretical analysis demonstrates that these algorithms achieve no-regret performance, even in the presence of adversarial costs and arbitrary delays. Extensive experiments on synthetic data validate the effectiveness of our approach, highlighting its potential in real-world applications where delayed feedback and nonsubmodular cost structures are prevalent.

**Keywords:** Nonsubmodular Minimization; Online Learning Algorithms; Delayed Feedback; No-Regret Learning; Bandit Feedback; Approximate Gradient Descent; Combinatorial Optimization

## 1 Introduction

In the rapidly evolving landscape of machine learning, the ability to make informed decisions in real-time is crucial. Many applications, from online advertising to autonomous systems, require algorithms that can adapt to uncertain environments while minimizing long-term costs. Traditionally, these problems have been addressed within the framework of submodular optimization, which benefits from well-established properties such as diminishing returns. However, many real-world challenges do not conform to submodular structures, instead involving complex nonsubmodular cost functions that defy these standard assumptions.

Nonsubmodular optimization presents unique difficulties, particularly in online settings where decisions must be made sequentially without complete information about future outcomes. The challenge is further compounded by the presence of delays in receiving feedback, a common occurrence in practical scenarios such as delayed rewards in reinforcement learning or lagging data in financial markets. These delays can significantly impair the performance of traditional online learning algorithms, which are typically designed under the assumption of immediate feedback.

This paper introduces a novel approach to online nonsubmodular minimization that addresses these challenges. We propose a suite of adaptive algorithms that leverage approximate gradient descent techniques, tailored specifically for nonsubmodular functions in both full-information and bandit feedback settings. To effectively manage delayed feedback, our algorithms incorporate a pooling strategy that dynamically adjusts to the timing of received information, ensuring robust performance even under significant delays.

**Backgrounds.** With machine learning systems increasingly being deployed in real-world settings, there is an urgent need for online learning algorithms that can minimize cumulative costs over the long run, even in the face of complete uncertainty about future outcomes. There exist a myriad of works that deal with this setting, most prominently in the area of online learning and bandits [Cesa-Bianchi and Lugosi(2006), Lattimore and Szepesvári(2020)]. The majority of this literature deals with problems where the decisions are taken from either a small set (such as in the multi armed bandit framework [Auer(2002)]), a continuous decision space (as in linear bandits [Auer(2002), Dani et al.(2008)]) or in the case the decision set is combinatorial in nature, the response is often assumed to maintain a simple functional relationship with the input (e.g., linear [Cesa-Bianchi and Lugosi(2012)]).

In this paper, we depart from these assumptions and explore what we believe is a more realistic type of model for the setting where the actions can be encoded as selecting a subset of a universe of size  $n$ . We study a sequential interaction between an agent and the world that takes place in rounds. At the beginning of round  $t$ , the agent chooses a *subset*  $S^t \subseteq [n]$  (e.g., selecting the set of products in a factory [McCormick(2005)]), after which the agent suffers cost  $f_t(S^t)$  such that  $f_t$  is an  $\alpha$ -weakly DR-submodular and  $\beta$ -weakly DL-supermodular function [Lehmann et al.(2006)]. The agent then may receive extra information about  $f_t$  as feedback, for example in the full information setting the agent observes the whole function  $f_t$  and in the bandit feedback scenario the learner does not receive any extra information about  $f_t$  beyond the value of  $f_t(S^t)$ . The standard metric to measure an online learning algorithm is *regret* [Blum and Mansour(2007)]: the regret at time  $T$  is the difference between  $\sum_{t=1}^T f_t(S^t)$  that is the total cost achieved by the algorithm and  $\min_{x \in A} \sum_{t=1}^T f_t(x)$  that is the total cost achieved by the best fixed action in hindsight. A *no-regret* learning algorithm is one that achieves sublinear regret (as a function of  $T$ ). Many no-regret learning algorithms have been developed based on online convex optimization toolbox [Zinkevich(2003), Kalai and Vempala(2005), Shalev-Shwartz and Singer(2006), Hazan et al.(2007), Shalev-Shwartz(2011), Arora et al.(2012), Hazan(2016)] many of them achieving minimax-optimal regret bounds for different cost functions even when these are produced by the world in an adversarial fashion. However, many online decision-making problems remain open, for example when the decision space is discrete and large (e.g., exponential in the number of problem parameters) and the cost functions are nonlinear [Hazan and Kale(2012)].

To the best of our knowledge, [Hazan and Kale(2012)] were the first to investigate non-parametric online learning in combinatorial domains by considering the setting where the costs  $f_t$  are all submodular functions. In this formulation the decision space is the set of all subsets of a set of  $n$  elements; and the cost functions are *submodular*. They provided no-regret algorithms for both the full information and bandit settings. Their chief innovation was to propose a computationally efficient algorithm for online submodular learning that resolved the exponential computational and statistical dependence on  $n$  suffered by all previous approaches [Hazan and Kale(2012)]. These results served as a catalyst for a rich and expanding research area [Streeter and Golovin(2008), Jegelka and Bilmes(2011), Buchbinder et al.(2014), Chen et al.(2018c), Roughgarden and Wang(2018), Chen et al.(2018b), Cardoso and Cummings(2019), Anari et al.(2019), Harvey et al.(2020), Thang and Srivastav(2021), Matsuoka et al.(2021)].

Even though submodularity can be used to model a few important typical cost functions that arise in machine learning problems [Boykov et al.(2001), Boykov and Kolmogorov(2004), Narasimhan et al.(2005), Bach(2010)], it is an insufficient assumption for many other applications where the cost functions do not satisfy submodularity, e.g., structured sparse learning [El Halabi and Cevher(2015)], batch Bayesian optimization [González et al.(2016), Bogunovic et al.(2016)], Bayesian A-optimal exper-

imental design [Bian et al.(2017)], column subset selection [Sviridenko et al.(2017)] and so on. In this work we aim to fill in this gap. In view of all this, we consider the following question: ***Can we design online learning algorithms when the cost functions are nonsubmodular?*** This paper provides an affirmative answer to this question by demonstrating that online/bandit approximate gradient descent algorithm can be directly extended from online submodular minimization [Hazan and Kale(2012)] to online nonsubmodular minimization when each cost functions  $f_t$  satisfy the regularity condition in [El Halabi and Jegelka(2020)].

Moreover, in online decision-making there is often a significant delay between decision and feedback. This delay has an adverse effect on the characterization between marketing feedback and an agent’s decision [Quanrud and Khashabi(2015), Héliou et al.(2020)]. For example, a click on an ad can be observed within seconds of the ad being displayed, but the corresponding sale can take hours or days to occur. We extend all of our algorithms to the delayed feedback setting by leveraging a pooling strategy recently introduced by [Héliou et al.(2020)] into the framework of online/bandit approximate gradient descent.

**Brief Contributions.** Our contributions are twofold: First, we extend the theory of online learning to accommodate the nonsubmodular setting, introducing a new regret measure that accounts for the inherent complexity of these functions. Second, we provide rigorous theoretical guarantees that our algorithms achieve sublinear regret, both in expectation and with high probability, across a range of delay scenarios.

## 1.1 Related Work

The offline nonsubmodular optimization with different notions of approximate submodularity has recently received a lot of attention. Most research focused on the maximization of nonsubmodular set functions, emerging as an important paradigm for studying real-world application problems [Das and Kempe(2011), Horel and Singer(2016), Chen et al.(2018a), Kuhnle et al.(2018), Hassidim and Singer(2018), Elenberg et al.(2018), Harshaw et al.(2019)]. In contrast, we are aware of relatively few investigations into the minimization of nonsubmodular set functions. An interesting example is the ratio problem [Bai et al.(2016)] where the objective function to be minimized is the ratio of two set functions and is thus nonsubmodular in general. Note that the ratio problem does not admit a constant factor approximation even when two set functions are submodular [Svitkina and Fleischer(2011)]. However, if the objective function to be minimized is approximately modular with bounded curvature, the optimal approximation algorithms exist even when the constrain sets are assumed [Iyer et al.(2013)]. Another typical example is the minimization of the difference of two submodular functions, where some approximation algorithms were proposed in [Iyer and Bilmes(2012)] and [Kawahara et al.(2015)] but without any approximation guarantee. Very recently, [El Halabi and Jegelka(2020)] provided a comprehensive treatment of optimal approximation guarantees for minimizing nonsubmodular set functions, characterized by how close the function is to submodular. Our work is close to theirs and our results can be interpreted as the extension of [El Halabi and Jegelka(2020)] to online learning with delayed feedback.

Another line of relevant works comes from online learning literature and focuses on no-regret algorithms in different settings with delayed costs. In the context of online convex optimization, [Quanrud and Khashabi(2015)] proposed an extension of online gradient descent (OGD) where the agent performs a batched gradient update the moment gradients are received and proved that OGD achieved a regret bound of  $O(\sqrt{T + D_T})$  where  $D_T$  is the total delay over a horizon  $T$ . However, their batch update approach can not be extended to bandit convex optimization since it does

not work with stochastic estimates of the received gradient information (or when attempting to infer such information from realized costs). This issue was posted by [Zhou et al.(2017)] and recently resolved by [Héliou et al.(2020)] who proposed a new pooling strategy based on a priority queue. The effect of delay was also discussed in the multi-armed bandit (MAB) literature under different assumptions [Joulani et al.(2013), Joulani et al.(2016), Vernade et al.(2017), Pike-Burke et al.(2018), Thune et al.(2019), Bistritz et al.(2019), Zhou et al.(2019), Zimmert and Seldin(2020), Gyorgy and Joulani(2021)]. In particular, [Thune et al.(2019)] proved the regret bound in adversarial MABs with the cumulative delay and [Gyorgy and Joulani(2021)] studied the adaptive tuning to delays and data in this setting. Further, [Joulani et al.(2016)] and [Zimmert and Seldin(2020)] also investigated adaptive tuning to the unknown sum of delays while [Bistritz et al.(2019)] and [Zhou et al.(2019)] gave further results in adversarial and linear contextual bandits respectively. However, the algorithms developed in the aforementioned works have little to do with online nonsubmodular minimization with delayed costs.

## 1.2 Contribution and Organization

First, we introduce a new notion of  $(\alpha, \beta)$ -regret which allows for analyzing no-regret online learning algorithms when the loss functions are nonsubmodular. We then propose two randomized algorithms for both the full-information and bandit feedback settings respectively with the regret bounds in expectation and high-probability sense. We then combine the aforementioned algorithms with the pooling strategy found in [Héliou et al.(2020)] and prove that the resulting algorithms are no-regret even when the delays are unbounded (cf. Assumption 1). Specifically, when the delay  $d_t$  satisfies  $d_t = o(t^\gamma)$ , we establish a  $O(\sqrt{nT^{1+\gamma}})$  regret bound in full-information setting and a  $O(nT^{\frac{2+\gamma}{3}})$  regret bound in bandit feedback setting. To our knowledge, this is the first theoretical guarantee for no-regret learning in online nonsubmodular minimization with delayed costs. Experimental results on sparse learning with synthetic data confirm our theoretical findings.

It is worth comparing our results with that in the existing works [El Halabi and Jegelka(2020), Hazan and Kale(2012), Héliou et al.(2020)]. First of all, the results concerning online nonsubmodular minimization are not a straightforward consequence of [El Halabi and Jegelka(2020)]. Indeed, it is natural yet nontrivial to identify the notion of  $(\alpha, \beta)$ -regret under which formal guarantees can be established for the nonsubmodular case. This notion does not appear before and appears to be a novel idea and an interesting conceptual contribution. Further, our results provide the first theoretical guarantee for no-regret learning in online and bandit nonsubmodular minimization and generalize the results in [Hazan and Kale(2012)]. Even though the online and bandit learning algorithms and regret analysis share the similar spirits with the context of [Hazan and Kale(2012)], the proof techniques are different since we need to deal with the nonsubmodular case with  $(\alpha, \beta)$ -regret. Finally, we are not aware of any results on online and bandit combinatorial optimization with delayed costs. [Héliou et al.(2020)] focused on the gradient-free game-theoretical learning with delayed costs where the action sets are *continuous* and *bounded*. Thus, their results can not imply ours. The only component that two works share is the pooling strategy which has been a common algorithmic component to handle the delays. Even though the pooling strategy is crucial to our delayed algorithms, we make much efforts to combine them properly and prove  $(\alpha, \beta)$ -regret bound of our new algorithms.

**Notation.** We let  $[n]$  be the set  $\{1, 2, \dots, n\}$  and  $\mathbb{R}_+^n$  be the set of all vectors in  $\mathbb{R}^n$  with non-negative components. We denote  $2^{[n]}$  as the set of all subsets of  $[n]$ . For a set  $S \subseteq [n]$ , we let  $\chi_S \in \{0, 1\}^n$  be the characteristic vector satisfying that  $\chi_S(i) = 1$  for each  $i \in S$  and  $\chi_S(i) = 0$

for each  $i \notin S$ . For a function  $f : 2^{[n]} \mapsto \mathbb{R}$ , we denote the marginal gain of adding an element  $i$  to  $S$  by  $f(i \mid S) = f(S \cup \{i\}) - f(S)$ . In addition,  $f$  is normalized if  $f(\emptyset) = 0$  and nondecreasing if  $f(A) \leq f(B)$  for  $A \subseteq B$ . For a vector  $x \in \mathbb{R}^n$ , its Euclidean norm refers to  $\|x\|$  and its  $i$ -th entry refers to  $x_i$ . We denote the support set of  $x$  by  $\text{supp}(x) = \{i \in [n] : x_i \neq 0\}$  and, by abuse of notation, we let  $x$  define a set function  $x(S) = \sum_{i \in S} x_i$ . We let  $P_S$  be the projection onto a closed set  $S$  and  $\text{dist}(x, S) = \inf_{y \in S} \|x - y\|$  denotes the distance between  $x$  and  $S$ . A pair of parameters  $(\alpha, \beta) \in \mathbb{R}_+ \times \mathbb{R}_+$  in the regret refer to approximation factors of the corresponding offline setting. Lastly,  $a = O(b(\alpha, \beta, n, T))$  refers to an upper bound  $a \leq C \cdot b(\alpha, \beta, n, T)$  where  $C > 0$  is independent of  $\alpha, \beta, n$  and  $T$ .

## 2 Preliminaries and Technical Background

We present the basic setup for minimizing structured nonsubmodular functions, including motivating examples and convex relaxation based on Lovász extension. We extend the offline setting to online setting and  $(\alpha, \beta)$ -regret which is important to the subsequent analysis.

### 2.1 Structured nonsubmodular function

Minimizing a set function  $f : 2^{[n]} \mapsto \mathbb{R}$  is NP-hard in general but is solved exactly with *submodular* structure in polynomial time [Iwata(2003), Grötschel et al.(2012), Lee et al.(2015)] and in strongly polynomial time [Schrijver(2000), Iwata et al.(2001), Iwata and Orlin(2009), Orlin(2009), Lee et al.(2015)]. More specifically,  $f$  is submodular if it satisfies the diminishing returns (DR) property as follows,

$$f(i \mid A) \geq f(i \mid B), \quad \text{for all } A \subseteq B, i \in [n] \setminus B \quad (1)$$

Further,  $f$  is modular if the inequality in Eq. (1) holds as an equality and is supermodular if

$$f(i \mid B) \geq f(i \mid A), \quad \text{for all } A \subseteq B, i \in [n] \setminus B$$

Relaxing these inequalities will bring us the notions of weak DR-submodularity/supermodularity that were introduced by [Lehmann et al.(2006)] and revisited in the machine learning literature [Bian et al.(2017)]. Formally, we have

**Definition 1.** A set function  $f : 2^{[n]} \mapsto \mathbb{R}$  is  $\alpha$ -weakly DR-submodular with  $\alpha > 0$  if

$$f(i \mid A) \geq \alpha f(i \mid B), \quad \text{for all } A \subseteq B, i \in [n] \setminus B$$

Similarly,  $f$  is  $\beta$ -weakly DR-supermodular with  $\beta > 0$  if

$$f(i \mid B) \geq \beta f(i \mid A), \quad \text{for all } A \subseteq B, i \in [n] \setminus B$$

We say that  $f$  is  $(\alpha, \beta)$ -weakly DR-modular if both of the above two inequalities hold true.

The above notions of weak DR-submodularity (or weak DR-supermodularity) generalize the notions of submodularity (or supermodularity); indeed, we have  $f$  is submodular (or supermodular) if and only if  $\alpha = 1$  (or  $\beta = 1$ ). They are also special cases of more general notions of weak submodularity (or weak supermodularity) [Das and Kempe(2011)] and we refer to [Bogunovic et al.(2018), Proposition 1] and [El Halabi et al.(2018), Proposition 8] for the details. For an overview of the approximate submodularity, we refer to [Bian et al.(2017), Section 6] and [El Halabi and Jegelka(2020),

Figure 1]. In addition, the parameters  $1 - \alpha$  and  $1 - \beta$  are referred to as *generalized inverse curvature* and *generalized curvature* respectively [Bian et al.(2017), Bogunovic et al.(2018)] and can be interpreted as the extension of inverse curvature and curvature [Conforti and Cornuéjols(1984)] for submodular and supermodular functions. Intuitively, these parameters quantify how far the function  $f$  is from being a submodular (or supermodular) function.

Recently, [El Halabi and Jegelka(2020)] have proposed and studied the problem of minimizing a class of structured nonsubmodular functions as follows,

$$\min_{S \subseteq [n]} f(S) := \bar{f}(S) - \underline{f}(S) \quad (2)$$

where  $\bar{f}$  and  $\underline{f}$  are both normalized (i.e.,  $\bar{f}(\emptyset) = \underline{f}(\emptyset) = 0$ )<sup>1</sup> and nondecreasing,  $\bar{f}$  is  $\alpha$ -weakly DR-submodular and  $\underline{f}$  is  $\beta$ -weakly DR-supermodular. Note that the problem in Eq. (2) is challenging; indeed,  $f$  is neither weakly DR-submodular nor weakly DR-supermodular in general since the weak DR-submodularity (or weak DR-supermodularity) are only valid for monotone functions.

It is worth mentioning that Eq. (2) is not necessarily theoretically artificial but encompasses a wide range of applications. We present two typical examples which can be formulated in the form of Eq. (2) and refer to [El Halabi and Jegelka(2020), Section 4] for more details.

**Example 1** (Structured Sparse Learning). We aim to estimate a sparse parameter vector whose support satisfies a particular structure and commonly formulate such problems as  $\min_{x \in \mathbb{R}^n} \ell(x) + \lambda f(\text{supp}(x))$ , where  $\ell : \mathbb{R}^n \mapsto \mathbb{R}$  is a loss function and  $f : 2^{[n]} \mapsto \mathbb{R}$  is a set function favoring the desirable supports. Existing approaches such as [Bach(2010)] proposed to replace the discrete regularization function  $f(\text{supp}(x))$  by its closest convex relaxation and is computationally tractable only when  $f$  is submodular. However, this problem is often better modeled by a nonsubmodular regularizer in practice [El Halabi and Cevher(2015)]. An alternative formulation of structured sparse learning problems is

$$\min_{S \subseteq [n]} \lambda f(S) - h(S) \quad (3)$$

where  $h(S) = \ell(0) - \min_{\text{supp}(x) \subseteq S} \ell(x)$ . Note that Eq. (3) can be reformulated into the form of Eq. (2) under certain conditions; indeed,  $h$  is a normalized and nondecreasing function and [El Halabi and Jegelka(2020), Proposition 5] has shown that  $h$  is weakly DR-modular if  $\ell$  is smooth, strongly convex and is generated from random data. Examples of weakly DR-submodular regularizers  $f$  include the ones used in time-series and cancer diagnosis [Rapaport et al.(2008)] and healthcare [Sakaue(2019)].

**Example 2** (Batch Bayesian Optimization). We aim to optimize an unknown expensive-to-evaluate noisy function  $\ell$  with as few batches of function evaluations as possible. The evaluation points are chosen to maximize an acquisition function – the variance reduction function [González et al.(2016)] – subject to a cardinality constraint. Maximizing the variance reduction may be phrased as a special instance of the problems in Eq. (2) in the form of  $\min_{S \subseteq [n]} \lambda |S| - G(S)$ , where  $G : 2^{[n]} \mapsto \mathbb{R}$  is the variance reduction function defined accordingly and [El Halabi and Jegelka(2020), Proposition 6] has shown that it is also non-decreasing and weakly DR-modular. This formulation allows to include nonlinear costs with (weak) decrease in marginal costs (economies of scale) with some applications in the sensor placement.

---

<sup>1</sup>In general, we can let  $\bar{f}(S) \leftarrow \bar{f}(S) - \bar{f}(\emptyset)$  and  $\underline{f}(S) \leftarrow \underline{f}(S) - \underline{f}(\emptyset)$  which will not change the minimization problem.

## 2.2 Convex relaxation based on the Lovász extension

The Lovász extension [Lovász(1983)] is a toolbox commonly used for minimizing a submodular set function  $f : 2^{[n]} \mapsto \mathbb{R}$ . It is a continuous interpolation of  $f$  on the unit hypercube  $[0, 1]^n$  and can be minimized efficiently since it is *convex* if and only if  $f$  is submodular. The minima of the Lovász extension also recover the minima of  $f$ .

Before the formal argument, we define a maximal chain of  $[n]$ ; that is,  $\{A_0, \dots, A_n\}$  is a maximal chain if  $\emptyset = A_0 \subseteq A_1 \subseteq \dots \subseteq A_n = [n]$ . Formally, we have

**Definition 2.** *Given a submodular function  $f$ , the Lovász extension is the function  $f_L : [0, 1]^n \mapsto \mathbb{R}$  given by  $f_L(x) = \sum_{i=0}^n \lambda_i f(A_i)$  where  $\{A_0, \dots, A_n\}$  is a maximal chain<sup>2</sup> of  $[n]$  so that  $\sum_{i=0}^n \lambda_i \chi_{A_i} = x$  and  $\sum_{i=0}^n \lambda_i = 1$  where  $\chi_{A_i}(j) = 1$  for  $\forall j \in A_i$  and  $\chi_{A_i}(j) = 0$  for  $\forall j \notin A_i$ .*

Even though Definition 2 implies that  $f_L(\chi_S) = f(S)$  for all  $S \subseteq [n]$ , it remains unclear how to find the chain or the coefficients. The preceding discussion defines the Lovász extension in an equivalent way that is more amenable for computing the subgradient of  $f_L$ .

Let  $x = (x_1, x_2, \dots, x_n) \in [0, 1]^n$  and we define that  $\pi : [n] \mapsto [n]$  is the sorting permutation of  $\{x_1, x_2, \dots, x_n\}$  where  $\pi(i) = j$  implies that  $x_j$  is the  $i$ -th largest element. By definition, we have  $1 \geq x_{\pi(1)} \geq \dots \geq x_{\pi(n)} \geq 0$  and let  $x_{\pi(0)} = 1$  and  $x_{\pi(n+1)} = 0$  for simplicity. Then, we set  $\lambda_i = x_{\pi(i)} - x_{\pi(i+1)}$  for all  $0 \leq i \leq n$  and let  $A_0 = \emptyset$  and  $A_i = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$ . We also have

$$\begin{aligned} \sum_{i=0}^n \lambda_i \chi_{A_i} &= \sum_{i=0}^n (x_{\pi(i)} - x_{\pi(i+1)}) (\chi_{A_{i-1}} + e_{\pi(i)}) \\ &= \sum_{i=1}^n e_{\pi(i)} \sum_{j=i}^n (x_{\pi(j)} - x_{\pi(j+1)}) = x \end{aligned}$$

As such, we obtain that  $f_L(x) = \sum_{i=1}^n x_{\pi(i)} f(\pi(i) \mid A_{i-1})$  where  $x_{\pi(1)} \geq x_{\pi(2)} \geq \dots \geq x_{\pi(n)}$  are the sorted entries in decreasing order,  $A_0 = \emptyset$  and  $A_i = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$ . Then, the classical results [Edmonds(2003), Fujishige(2005)] suggest that the subgradient  $g$  of  $f_L$  at any  $x \in [0, 1]^n$  can be computed by simply sorting the entries in decreasing order and taking

$$g_{\pi(i)} = f(A_i) - f(A_{i-1}), \text{ for all } i \in [n] \quad (4)$$

Since  $f_L$  is convex if and only if  $f$  is submodular, we can apply the convex optimization toolbox here. Recently, [El Halabi and Jegelka(2020)] have shown that the similar idea can be extended to nonsubmodular optimization in Eq. (2).

More specifically, we can define the convex closure  $f_C$  for any nonsubmodular function  $f$ ; indeed,  $f_C : [0, 1]^n \mapsto \mathbb{R}$  is the point-wise largest convex function which always lower bounds  $f$ . By definition,  $f_C$  is the *tightest* convex extension of  $f$  and  $\min_{S \subseteq [n]} f(S) = \min_{x \in [0, 1]^n} f_C(x)$ . In general, it is NP-hard to evaluate and optimize  $f_C$  [Vondrák(2007)]. Fortunately, [El Halabi and Jegelka(2020)] demonstrated that the Lovász extension  $f_L$  approximates  $f_C$  such that the vector computed using the approach in [Edmonds(2003)] and [Fujishige(2005)] approximates the subgradient of  $f_C$ . We summarize their results in the following proposition and provide the proofs in Section 2.4 for completeness.

---

<sup>2</sup>Both the chain and the set of  $\lambda_i$  may depend on the input  $x$ .

**Proposition 1.** *Focusing on Eq. (2), we let  $x \in [0, 1]^n$  with  $x_{\pi(1)} \geq \dots \geq x_{\pi(n)}$  and  $g_{\pi(i)} = f(A_i) - f(A_{i-1})$  for all  $i \in [n]$  where  $A_0 = \emptyset$  and  $A_i = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$ . Then, we have*

$$f_L(x) = g^\top x \geq f_C(x) \quad (5)$$

and

$$g(A) = \sum_{i \in A} g_i \leq \frac{1}{\alpha} \bar{f}(A) - \beta \underline{f}(A), \text{ for all } A \subseteq [n] \quad (6)$$

and

$$g^\top z \leq \frac{1}{\alpha} \bar{f}_C(z) + \beta(-\underline{f})_C(z), \text{ for all } z \in [0, 1]^n \quad (7)$$

Proposition 1 highlights how  $f_L$  approximates  $f_C$ ; indeed, we see from Eq. (5) and Eq. (7) that  $f_C(x) \leq f_L(x) \leq \frac{1}{\alpha} \bar{f}_C(x) + \beta(-\underline{f})_C(x)$  for all  $x \in [0, 1]^n$ . As such, it gives the key insight for analyzing the offline algorithms in [El Halabi and Jegelka(2020)] and will play an important role in the subsequent analysis of our paper.

### 2.3 Online nonsubmodular minimization

We consider online nonsubmodular minimization which extends the offline problem in Eq. (2) to the online setting. In particular, an adversary first chooses structured nonsubmodular functions  $f_1, f_2, \dots, f_T : 2^{[n]} \mapsto \mathbb{R}$  given by

$$f_t(S) := \bar{f}_t(S) - \underline{f}_t(S), \text{ for all } S \subseteq [n], t \in [T] \quad (8)$$

where  $\bar{f}_t$  and  $\underline{f}_t$  are normalized and non-decreasing,  $\bar{f}_t$  is  $\alpha$ -weakly DR-submodular and  $\underline{f}_t$  is  $\beta$ -weakly DR-supermodular. In each round  $t = 1, 2, \dots, T$ , the agent chooses  $S^t$  and observes the incurred loss  $f_t(S^t)$  after committing to her decision. Throughout the horizon  $[0, T]$ , one aims to minimize the regret – the difference between  $\sum_{t=1}^T f_t(S^t)$  and the loss at the best fixed solution in hindsight, i.e.,  $S_\star^T = \operatorname{argmin}_{S \subseteq [n]} \sum_{t=1}^T f_t(S)$  – which is defined by<sup>3</sup>

$$R(T) = \sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T f_t(S_\star^T) \quad (9)$$

An algorithm is *no-regret* if  $R(T)/T \rightarrow 0$  as  $T \rightarrow +\infty$  and *efficient* if it computes each decision set  $S^t$  in polynomial time. In the context, the regret is used when the minimization for a known cost, i.e.,  $\min_{S \subseteq [n]} f(S)$ , can be solved exactly. However, solving the optimization problem in Eq. (2) with nonsubmodular costs is NP-hard regardless of any multiplicative constant factor [Iyer and Bilmes(2012), Trevisan(2014)]. Thus, it is necessary to consider a bicriteria-like approximation guarantee with the factors  $\alpha, \beta > 0$  as [El Halabi and Jegelka(2020)] suggested. In particular,  $(\alpha, \beta)$  are bounds on the quality of a solution  $S$  returned by a given offline algorithm compared to the optimal solution  $S_\star$ ; that is,  $f(S) \leq \frac{1}{\alpha} \bar{f}(S_\star) - \beta \underline{f}(S_\star)$ . Such bicriteria-like approximation is optimal: [El Halabi and Jegelka(2020), Theorem 2] has shown that no algorithm with subexponential number of value queries can improve on it in the oracle model.

Our goal is to analyze *online approximate gradient descent algorithm and its bandit variant* for online nonsubmodular minimization. Let  $(\alpha, \beta)$  be the approximation factors attained by an offline algorithm that solves  $\min_{S \subseteq [n]} f(S)$  for a known nonsubmodular function  $f$  in Eq. (2). The

<sup>3</sup>If the sets  $S^t$  are chosen by a randomized algorithm, we consider the expected regret over the randomness.



$(\alpha, \beta)$ -*regret* compares to the best solution that can be expected in polynomial time and is defined by

$$R_{\alpha, \beta}(T) = \sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T \left( \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T) \right) \quad (10)$$

where  $S_\star^T = \operatorname{argmin}_{S \subseteq [n]} \sum_{t=1}^T f_t(S)$ . It is analogous to the  $\alpha$ -regret which is widely used in online constrained submodular minimization [Jegelka and Bilmes(2011)] and online submodular maximization [Streeter and Golovin(2008)].

As mentioned before, we consider the algorithmic design in both *full information* and *bandit feedback* settings. In the former one, the agent is allowed to have unlimited access to the value oracles of  $f_t(\cdot)$  after choosing  $S^t$  in each round  $t$ . In the latter one, the agent only observes the incurred loss at the point that she has chosen in each round  $t$ , i.e.,  $f_t(S^t)$ , and receives no other information.

## 2.4 Proof of Proposition 1

We have

$$f_C(x) = \max_{g \in \mathbb{R}^d, \rho \in \mathbb{R}} \left\{ g^\top x + \rho : g(A) + \rho \leq f(A), \forall A \subseteq [n] \right\} \quad (11)$$

First, we prove Eq. (5) using the definition of  $f_L$  and Eq. (11). Indeed, we have  $f_L(x) = g^\top x$  where we let  $x \in [0, 1]^n$  with  $x_{\pi(1)} \geq \dots \geq x_{\pi(n)}$  and  $g_{\pi(i)} = f(\pi(i) \mid A_{i-1})$  for all  $i \in [n]$ . Then, it suffices to show that  $g^\top x \geq \tilde{g}^\top x + \tilde{\rho}$  in which  $\tilde{g}(A) + \tilde{\rho} \leq f(A)$  for all  $A \subseteq [n]$ . We have

$$\begin{aligned} g^\top x - (\tilde{g}^\top x + \tilde{\rho}) &= \sum_{i=1}^n x_{\pi(i)} (f(\pi(i) \mid A_{i-1}) - \tilde{g}_{\pi(i)}) - \tilde{\rho} \\ &= \sum_{i=1}^{n-1} (x_{\pi(i)} - x_{\pi(i+1)}) (f(A_i) - \tilde{g}(A_i)) + x_{\pi(n)} (f([n]) - \tilde{g}([n])) - \tilde{\rho} \end{aligned}$$

Since  $\tilde{g}(A) + \tilde{\rho} \leq f(A)$  for all  $A \subseteq [n]$ , we have

$$f([n]) - \tilde{g}([n]) \geq \tilde{\rho}, \quad f(A_i) - \tilde{g}(A_i) \geq \tilde{\rho}, \quad \text{for all } i \in [n]$$

Putting these pieces together with  $x_{\pi(1)} \geq \dots \geq x_{\pi(n)}$  yields that

$$g^\top x - (\tilde{g}^\top x + \tilde{\rho}) \geq \sum_{i=1}^{n-1} (x_{\pi(i)} - x_{\pi(i+1)}) \tilde{\rho} + x_{\pi(n)} \tilde{\rho} - \tilde{\rho} = (x_{\pi(1)} - 1) \tilde{\rho}$$

Since  $x \in [0, 1]^n$ , we have  $x_{\pi(1)} \leq 1$ . Since  $\tilde{g}(A) + \tilde{\rho} \leq f(A)$  for all  $A \subseteq [n]$  and  $f(\emptyset) = 0$ , we derive by letting  $A = \emptyset$  that  $\tilde{\rho} \leq f(\emptyset) - \tilde{g}(\emptyset) \leq 0$ . This implies the desired result.

Further, we prove Eq. (6) using the definition of weak DR-submodularity. Indeed, we have  $g(A) = \sum_{i \in A} g_i$ . Since  $g_{\pi(i)} = f(\pi(i) \mid A_{i-1})$  for all  $i \in [n]$ , we have

$$g(A) = \sum_{\pi(i) \in A} f(\pi(i) \mid A_{i-1}) = \sum_{\pi(i) \in A} (\bar{f}(\pi(i) \mid A_{i-1}) - \underline{f}(\pi(i) \mid A_{i-1}))$$

Since  $\bar{f}$  is  $\alpha$ -weakly DR-submodular,  $\underline{f}$  is  $\beta$ -weakly DR-supermodular and  $A \cap A_{i-1} \subseteq A_{i-1}$ , we have

$$\bar{f}(\pi(i) \mid A \cap A_{i-1}) \geq \alpha \bar{f}(\pi(i) \mid A_{i-1}), \quad \underline{f}(\pi(i) \mid A_{i-1}) \geq \beta \underline{f}(\pi(i) \mid A \cap A_{i-1}) \quad (12)$$

Putting these pieces together yields that

$$g(A) \leq \sum_{\pi(i) \in A} \left( \frac{1}{\alpha} \bar{f}(\pi(i) \mid A \cap A_{i-1}) - \beta \underline{f}(\pi(i) \mid A \cap A_{i-1}) \right)$$

Then, we have

$$\begin{aligned} g(A) &\leq \frac{1}{\alpha} \left( \sum_{i=1}^n (\bar{f}(A \cap A_i) - \bar{f}(A \cap A_{i-1})) \right) - \beta \left( \sum_{i=1}^n (\underline{f}(A \cap A_i) - \underline{f}(A \cap A_{i-1})) \right) \\ &= \frac{1}{\alpha} \bar{f}(A) - \beta \underline{f}(A), \quad \text{for all } A \subseteq [n] \end{aligned}$$

This implies the desired result.

Finally, we prove Eq. (7) using Eq. (11). Indeed, we have  $g = \bar{g} - \underline{g}$  where  $\bar{g}_{\pi(i)} = \bar{f}(\pi(i) \mid A_{i-1})$  and  $\underline{g}_{\pi(i)} = \underline{f}(\pi(i) \mid A_{i-1})$  for all  $i \in [n]$ . For any  $A \subseteq [n]$ , we obtain by using Eq. (12) that

$$\begin{aligned} \bar{g}(A) &\leq \frac{1}{\alpha} \left( \sum_{\pi(i) \in A} \bar{f}(\pi(i) \mid A \cap A_{i-1}) \right) = \frac{1}{\alpha} \left( \sum_{i=1}^n (\bar{f}(A \cap A_i) - \bar{f}(A \cap A_{i-1})) \right) = \frac{1}{\alpha} \bar{f}(A) \\ -\underline{g}(A) &\leq -\beta \left( \sum_{\pi(i) \in A} \underline{f}(\pi(i) \mid A \cap A_{i-1}) \right) = -\beta \left( \sum_{i=1}^n (\underline{f}(A \cap A_i) - \underline{f}(A \cap A_{i-1})) \right) = -\beta \underline{f}(A). \end{aligned}$$

Equivalently, we have  $\alpha \bar{g}(A) + 0 \leq \bar{f}(A)$  and  $\frac{1}{\beta}(-\underline{g}(A)) + 0 \leq -\underline{f}(A)$  for any  $A \subseteq [n]$ . Using Eq. (11), we have

$$\alpha \bar{g}^\top z + 0 \leq \bar{f}_C(z), \quad \frac{1}{\beta}(-\underline{g})^\top z + 0 \leq (-\underline{f})_C(z), \quad \text{for all } z \in [0, 1]^n$$

Since  $g = \bar{g} - \underline{g}$  and  $\alpha, \beta > 0$ , we have  $g^\top z \leq \frac{1}{\alpha} \bar{f}_C(z) + \beta(-\underline{f})_C(z)$ . This implies the desired result.

### 3 Online Approximation Algorithm: Full Information Setting

In this and next sections, we analyze online approximate gradient descent algorithm and its bandit variant for regret minimization when the nonsubmodular cost functions are in the form of Eq (8). Here we discuss the full information setting.

Let  $[0, 1]^n$  be the unit hypercube and the cost function on  $[0, 1]^n$  corresponding to  $f_t$  is the function  $(f_t)_C$  that is the convex closure of  $f_t$ . Equipped with Proposition 1, we can compute approximate subgradients of  $(f_t)_C$  such that the online gradient descent [Zinkevich(2003)] is applicable.

This leads to Algorithm 1 which performs one-step projected gradient descent that yields  $x^t$  and then samples  $S^t$  from the distribution  $\lambda$  over  $\{A_i\}_{i=0}^n$  encoded by  $x^t$ . It is worth mentioning that  $\lambda_i^t = x_{\pi(i)}^t - x_{\pi(i+1)}^t$  for all  $0 \leq i \leq n$  and  $\lambda$  is thus completely independent of  $f_t$ . This guarantees that Algorithm 1 is valid in online manner since  $f_t$  is realized after the decision maker chooses  $S^t$ . One of the advantages of Algorithm 1 is that it does not require the value of  $\alpha$  and  $\beta$  which can be hard to compute in practice. We summarize our results for Algorithm 1 in the following theorem.

---

**Algorithm 1** Online Approximate Gradient Descent

---

- 1: **Initialization:** the point  $x^1 \in [0, 1]^n$  and the stepsize  $\eta > 0$
  - 2: **for**  $t = 1, 2, \dots$  **do**
  - 3:   Let  $x_{\pi(1)}^t \geq \dots x_{\pi(n)}^t$  be the sorted entries in the decreasing order with  $A_i^t = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$  and  $A_0^t = \emptyset$ . Let  $x_{\pi(0)}^t = 1$  and  $x_{\pi(n+1)}^t = 0$
  - 4:   Let  $\lambda_i^t = x_{\pi(i)}^t - x_{\pi(i+1)}^t$  for all  $0 \leq i \leq n$
  - 5:   Sample  $S^t$  from the distribution  $\mathbb{P}(S^t = A_i^t) = \lambda_i^t$  for all  $0 \leq i \leq n$  and observe the new loss function  $f_t$
  - 6:   Compute  $g_{\pi(i)}^t = f_t(A_i^t) - f_t(A_{i-1}^t)$  for all  $i \in [n]$
  - 7:   Compute  $x^{t+1} = P_{[0,1]^n}(x^t - \eta g^t)$
  - 8: **end for**
- 

**Theorem 1.** Suppose the adversary chooses nonsubmodular functions in Eq. (8) satisfying  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$ . Fixing  $T \geq 1$  and letting  $\eta = \frac{\sqrt{n}}{L\sqrt{T}}$  in Algorithm 1, we have  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(\sqrt{nT})$  and  $R_{\alpha,\beta}(T) = O(\sqrt{nT} + \sqrt{T \log \frac{1}{\delta}})$  with probability  $1 - \delta$ .

**Remark.** Theorem 1 demonstrates that Algorithm 1 is regret-optimal for our setting; indeed, our setting includes online unconstrained submodular minimization as a special case where  $(\alpha, \beta)$ -regret becomes standard regret in Eq. (9) and [Hazan and Kale(2012)] shows that Algorithm 1 is optimal up to constants. Our theoretical result also extends the results in [Hazan and Kale(2012)] from submodular cost functions to nonsubmodular cost functions in Eq. (8) using the  $(\alpha, \beta)$ -regret instead of the standard regret in Eq. (9).

### 3.1 Regret Analysis for Algorithm 1

In this section, we present several technical lemmas for analyzing the regret minimization property of Algorithm 1. We also give the missing proof of Theorem 1.

We provide two technical lemmas for Algorithm 1. The first lemma gives a bound on the vector  $g^t$  and the difference between  $x^t$  and any fixed  $x \in [0, 1]^n$ .

**Lemma 1.** Suppose that the iterates  $\{x^t\}_{t \geq 1}$  and the vectors  $\{g^t\}_{t \geq 1}$  be generated by Algorithm 1 and  $x \in [0, 1]^n$  and let  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$  and both  $\bar{f}_t$  and  $\underline{f}_t$  are nondecreasing. Then, we have  $\|x^t - x\| \leq \sqrt{n}$  and  $\|g^t\| \leq L$  for all  $t \geq 1$ .

The second lemma gives a key inequality for analyzing Algorithm 1.

**Lemma 2.** Suppose that the iterates  $\{x^t\}_{t \geq 1}$  are generated by Algorithm 1 and  $x \in [0, 1]^n$  and let  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ . Then, we have

$$\sum_{t=1}^T \mathbb{E}[(f_t)_L(x^t)] \leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) + \frac{n}{2\eta} + \frac{\eta L^2 T}{2}$$

*Proof of Theorem 1.* By the definition of the Lovász extension, we have

$$(f_t)_L(x^t) = \sum_{i=1}^{n-1} (x_{\pi(i)}^t - x_{\pi(i+1)}^t) f_t(A_i^t) + (1 - x_{\pi(1)}^t) f_t(A_0^t) + x_{\pi(n)}^t f_t(A_n^t) = \sum_{i=0}^n \lambda_i^t f_t(A_i^t)$$

By the update formula, we have  $\mathbb{E}[f_t(S^t) \mid x^t] = (f_t)_L(x^t)$  which implies that  $\mathbb{E}[f_t(S^t)] = \mathbb{E}[(f_t)_L(x^t)]$ . By the definition of the convex closure, we obtain that the convex closure of a set function  $f$  agrees with  $f$  on all the integer points [Dughmi(2009), Page 4, Proposition 3.3]. Letting  $S_\star^T = \operatorname{argmin}_{S \subseteq [n]} \sum_{t=1}^T f_t(S)$ , we have  $S_\star^T$  is an integer point and

$$(\bar{f}_t)_C(\chi_{S_\star^T}) = \bar{f}_t(S_\star^T), \quad (-\underline{f}_t)_C(\chi_{S_\star^T}) = -\beta \underline{f}_t(S_\star^T)$$

which implies that

$$\frac{1}{\alpha} (\bar{f}_t)_C(\chi_{S_\star^T}) + \beta (-\underline{f}_t)_C(\chi_{S_\star^T}) = \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T)$$

Putting these pieces together and letting  $x = \chi_{S_\star^T}$  in the inequality of Lemma 2 yields that

$$\sum_{t=1}^T \mathbb{E}[f_t(S^t)] \leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T) \right) + \frac{n}{2\eta} + \frac{\eta L^2 T}{2}$$

Plugging the choice of  $\eta = \frac{\sqrt{n}}{L\sqrt{T}}$  into the above inequality yields that  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(\sqrt{nT})$  as desired.

We proceed to derive a high probability bound using the concentration inequality. In particular, we review the Hoeffding inequality [Hoeffding(1963)] and refer to [Cesa-Bianchi and Lugosi(2006), Appendix A] for a proof. The following proposition is a restatement of [Cesa-Bianchi and Lugosi(2006), Corollary A.1].

**Proposition 2.** *Let  $X_1, \dots, X_n$  be independent real-valued random variables such that for each  $i = 1, \dots, n$ , there exist some  $a_i \leq b_i$  such that  $\mathbb{P}(a_i \leq X_i \leq b_i) = 1$ . Then for every  $\epsilon > 0$ , we have*

$$\begin{aligned} \mathbb{P}\left(\sum_{i=1}^n X_i - \mathbb{E}\left[\sum_{i=1}^n X_i\right] > +\epsilon\right) &\leq \exp\left(-\frac{2\epsilon^2}{\sum_{i=1}^n (b_i - a_i)^2}\right) \\ \mathbb{P}\left(\sum_{i=1}^n X_i - \mathbb{E}\left[\sum_{i=1}^n X_i\right] < -\epsilon\right) &\leq \exp\left(-\frac{2\epsilon^2}{\sum_{i=1}^n (b_i - a_i)^2}\right) \end{aligned}$$

Since the sequence of points  $x^1, x^2, \dots, x^T$  is obtained by several deterministic gradient descent steps, we have this sequence is purely deterministic. Since each of  $S^t$  is obtained by independent randomized rounding on the point  $x^t$ , we have the sequence of random variables  $X_t = f_t(S^t)$  is independent. By definition of  $f_t$ , we have

$$|X_t| = |\bar{f}_t(S^t) - \underline{f}_t(S^t)| \leq \bar{f}_t(S^t) + \underline{f}_t(S^t)$$

Since  $\bar{f}_t$  and  $\underline{f}_t$  are non-decreasing and  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ , we have  $\mathbb{P}(-L \leq X_i \leq L) = 1$  for all  $t \geq 1$ . Then, by Proposition 2, we have

$$\mathbb{P}\left(\sum_{i=1}^n f_t(S^t) - \mathbb{E}\left[\sum_{i=1}^n f_t(S^t)\right] > \epsilon\right) \leq \exp\left(-\frac{\epsilon^2}{2nL^2}\right)$$

Equivalently, we have  $\sum_{i=1}^n f_t(S^t) - \mathbb{E}[\sum_{i=1}^n f_t(S^t)] \leq L\sqrt{2T \log \frac{1}{\delta}}$  with probability at least  $1 - \delta$ . This together with  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(\sqrt{nT})$  yields that  $R_{\alpha,\beta}(T) = O(\sqrt{nT} + \sqrt{T \log \frac{1}{\delta}})$  with probability at least  $1 - \delta$  as desired.  $\square$

---

**Algorithm 2** Bandit Approximate Gradient Descent

---

- 1: **Initialization:** the point  $x^1 \in [0, 1]^n$  and the stepsize  $\eta > 0$ ; the exploration probability  $\mu \in (0, 1)$
  - 2: **for**  $t = 1, 2, \dots, T$  **do**
  - 3:   Let  $x_{\pi(1)}^t \geq \dots x_{\pi(n)}^t$  be the sorted entries in decreasing order with  $A_i^t = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$  and  $A_0^t = \emptyset$ . Let  $x_{\pi(0)}^t = 1$  and  $x_{\pi(n+1)}^t = 0$
  - 4:   Let  $\lambda_i^t = x_{\pi(i)}^t - x_{\pi(i+1)}^t$  for all  $0 \leq i \leq n$
  - 5:   Sample  $S^t$  from the distribution  $\mathbb{P}(S^t = A_i^t) = (1 - \mu)\lambda_i^t + \frac{\mu}{n+1}$  for all  $0 \leq i \leq n$  and observe the loss  $f_t(S^t)$
  - 6:   Compute  $\hat{f}_i^t = \frac{\mathbf{1}(S^t = A_i^t)}{(1 - \mu)\lambda_i^t + \mu/(n+1)} f_t(S^t)$  for all  $0 \leq i \leq n$
  - 7:   Compute  $\hat{g}_{\pi(i)}^t = \hat{f}_i^t - \hat{f}_{i-1}^t$  for all  $i \in [n]$
  - 8:   Compute  $x^{t+1} = P_{[0, 1]^n}(x^t - \eta g^t)$
  - 9: **end for**
- 

## 4 Online Approximation Algorithm: Bandit Feedback Setting

In contrast with the full-information setting, the agent only observes the loss function  $f_t$  at her action  $S^t$ , i.e.,  $f_t(S^t)$ , in bandit feedback setting. This is a more challenging setup since the agent does not have full access to the new loss function  $f_t$  at each round  $t$  yet.

Despite the bandit feedback, we can compute an unbiased estimator of the gradient  $g^t$  in Algorithm 1 using the technique of importance weighting and try to implement a stochastic version of Algorithm 1. More specifically, we notice that  $\hat{f}_i^t = \frac{\mathbf{1}(S^t = A_i^t)}{\lambda_i^t} f_t(S^t)$  is unbiased for estimating  $f_t(A_i^t)$  for all  $0 \leq i \leq n$ . Thus,  $\hat{g}_{\pi(i)}^t = \hat{f}_i^t - \hat{f}_{i-1}^t$  for all  $i \in [n]$  gives us an unbiased estimator of the gradient  $g^t$ . However, the variance of the estimator  $\hat{g}$  could be undesirably large since the values of  $\lambda_i^t$  may be arbitrarily small.

To resolve this issue, we can sample  $S^t$  from a mixture distribution that combines (with probability  $1 - \mu$ ) samples from  $\lambda^t$  and (with probability  $\mu$ ) samples from the uniform distribution over  $\{A_i^t\}_{i=0}^n$ . This guarantees that the variance of  $\hat{f}_i^t$  is upper bounded by  $O(n^2/\mu)$ . The similar idea has been employed in [Hazan and Kale(2012)] for online submodular minimization. Then, we conduct the careful analysis for the estimators  $\hat{g}^t$  such that the scale of the variance is taken into account. Note that our analysis is different from the standard analysis in [Flaxman et al.(2005)] which seems oversimplified for our setting and results in worse regret of  $O(T^{3/4})$  compared to our result in the following theorem.

**Theorem 2.** *Suppose the adversary chooses nonsubmodular functions  $f_t$  in Eq. (8) satisfying  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$ . Fixing  $T \geq 1$  and letting  $(\eta, \mu) = (\frac{1}{LT^{2/3}}, \frac{n}{T^{1/3}})$  in Algorithm 2, we have  $\mathbb{E}[R_{\alpha, \beta}(T)] = O(nT^{\frac{2}{3}})$  and  $R_{\alpha, \beta}(T) = O(nT^{\frac{2}{3}} + \sqrt{n \log \frac{1}{\delta}} T^{\frac{2}{3}})$  with probability  $1 - \delta$ .*

**Remark.** Theorem 2 demonstrates that Algorithm 2 is no-regret for our setting even when only the bandit feedback is available, further extending the results in [Hazan and Kale(2012)] from submodular cost functions to nonsubmodular cost functions in Eq. (8) using the  $(\alpha, \beta)$ -regret instead of the standard regret in Eq. (9).

## 4.1 Regret Analysis for Algorithm 2

In this section, we present several technical lemmas for analyzing the regret minimization property of Algorithm 2. We also give the missing proofs of Theorem 2.

We provide several technical lemmas for Algorithm 2. The first lemma is analogous to Lemma 1 and gives a bound on the vector  $\hat{g}^t$  (in expectation) and the difference between  $x^t$  and any fixed  $x \in [0, 1]^n$ .

**Lemma 3.** *Suppose that the iterates  $\{x^t\}_{t \geq 1}$  and the vectors  $\{\hat{g}^t\}_{t \geq 1}$  be generated by Algorithm 2 and  $x \in [0, 1]^n$  and let  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$  and both  $\bar{f}_t$  and  $\underline{f}_t$  are nondecreasing. Then, we have  $\|x^t - x\| \leq \sqrt{n}$  for all  $t \geq 1$  and*

$$\mathbb{E}[\hat{g}^t \mid x^t] = g^t, \quad \mathbb{E}[\|\hat{g}^t\|^2 \mid x^t] \leq \frac{8n^2 L^2}{\mu}, \quad \|\hat{g}^t\|^2 \leq \frac{2(n+1)^2 L^2}{\mu^2}$$

where we have  $g_{\pi(i)}^t = f_t(A_i^t) - f_t(A_{i-1}^t)$  for all  $i \in [n]$ .

The second lemma is analogous to Lemma 2 and gives a key inequality for analyzing Algorithm 2.

**Lemma 4.** *Suppose that the iterates  $\{x^t\}_{t \geq 1}$  are generated by Algorithm 2 and  $x \in [0, 1]^n$  and let  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ . Then, we have*

$$\sum_{t=1}^T \mathbb{E}[(f_t)_L(x^t)] \leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) + \frac{n}{2\eta} + \frac{4n^2 L^2 \eta T}{\mu}$$

To prove the high probability bound, we require the following concentration inequality. In particular, we review the Bernstein inequality for martingales [Freedman(1975)] and refer to [Cesa-Bianchi and Lugosi(2006), Appendix A] for a proof. The following proposition is a consequence of [Cesa-Bianchi and Lugosi(2006), Lemma A.8].

**Proposition 3.** *Let  $X_1, \dots, X_n$  be a bounded martingale difference sequence with respect to the filtration  $\mathcal{F} = (\mathcal{F}_i)_{1 \leq i \leq n}$  such that  $|X_i| \leq K$  for each  $i = 1, \dots, n$ . We also assume that  $\mathbb{E}[\|X_{i+1}\|^2 \mid \mathcal{F}_i] \leq V$  for each  $i = 1, \dots, n-1$ . Then for every  $\delta > 0$ , we have*

$$\mathbb{P} \left( \left| \sum_{i=1}^n X_i - \mathbb{E}[X_i \mid \mathcal{F}_{i-1}] \right| > \sqrt{2TV \log \frac{1}{\delta}} + \frac{\sqrt{2}}{3} K \log \frac{1}{\delta} \right) \leq \delta$$

Then we provide our last lemma which significantly generalizes Lemma 4 for deriving the high-probability bounds.

**Lemma 5.** *Suppose that the iterates  $\{x^t\}_{t \geq 1}$  are generated by Algorithm 2 with  $\mu = \frac{n}{T^{1/3}}$  and  $x \in [0, 1]^n$  and let  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ . Fixing a sufficiently small  $\delta \in (0, 1)$  and letting  $T > \log^{\frac{3}{2}}(1/\delta)$ . Then, we have*

$$\sum_{t=1}^T (f_t)_L(x^t) \leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{n}{2\eta} + \frac{4n^2 L^2 \eta T}{\mu} + 12LT^{\frac{2}{3}} \sqrt{n^2 + n \log \frac{1}{\delta}} + 6\eta L^2 T \sqrt{n \log \frac{1}{\delta}}$$

with probability at least  $1 - 3\delta$ .

*Proof of Theorem 2.* By the definition of the Lovász extension and  $\lambda^t$ , we have

$$(f_t)_L(x^t) = \sum_{i=1}^{n-1} (x_{\pi(i)}^t - x_{\pi(i+1)}^t) f_t(A_i^t) + (1 - x_{\pi(1)}^t) f_t(A_0^t) + x_{\pi(n)}^t f_t(A_n^t) = \sum_{i=0}^n \lambda_i^t f_t(A_i^t)$$

By the update formula of  $S^t$ , we have

$$\mathbb{E}[f_t(S^t) \mid x^t] - (f_t)_L(x^t) = \mu \sum_{i=0}^n \left( \frac{1}{n+1} - \lambda_i^t \right) f_t(A_i^t) \leq \mu \sum_{i=0}^n \left( \frac{1}{n+1} + \lambda_i^t \right) |f_t(A_i^t)|$$

Since  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$  and  $\bar{f}_t$  and  $\underline{f}_t$  are both normalized and non-decreasing, we have

$$\mathbb{E}[f_t(S^t) \mid x^t] - (f_t)_L(x^t) \leq L\mu \sum_{i=0}^n \left( \frac{1}{n+1} + \lambda_i^t \right) = 2L\mu \quad (13)$$

which implies that

$$\mathbb{E}[f_t(S^t)] - \mathbb{E}[(f_t)_L(x^t)] \leq 2L\mu$$

Using the same argument as in Theorem 1, we have

$$\frac{1}{\alpha}(\bar{f}_t)_C(\chi_{S_\star^T}) + \beta(-\underline{f}_t)_C(\chi_{S_\star^T}) = \frac{1}{\alpha}f_t(S_\star^T) - \beta\underline{f}_t(S_\star^T), \quad \text{where } S_\star^T = \operatorname{argmin}_{S \subseteq [n]} \sum_{t=1}^T f_t(S)$$

Putting these pieces together and letting  $x = \chi_{S_\star^T}$  in the inequality of Lemma 4 yields that

$$\sum_{t=1}^T \mathbb{E}[f_t(S^t)] \leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T) \right) + \frac{n}{2\eta} + \frac{4n^2 L^2 \eta T}{\mu} + 2LT\mu$$

Plugging the choice of  $\eta = \frac{1}{LT^{2/3}}$  and  $\mu = \frac{n}{T^{1/3}}$  into the above inequality yields that  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(nT^{\frac{2}{3}})$  as desired.

We proceed to derive a high probability bound using Lemma 5. Indeed, we first consider the case of  $T < 2 \log^{\frac{3}{2}}(1/\delta)$ . Since  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ , we have

$$R_{\alpha,\beta}(T) \leq \sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T \left( \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T) \right) \leq \left( 1 + \frac{1}{\alpha} + \beta \right) LT = O(T^{\frac{2}{3}} \sqrt{\log \frac{1}{\delta}})$$

For the case of  $T \geq 2 \log^{\frac{3}{2}}(1/\delta)$ , we obtain by combining Lemma 5 with Eq. (13) that

$$\begin{aligned} \sum_{t=1}^T \mathbb{E}[f_t(S^t) \mid x^t] &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) \\ &\quad + \frac{n}{2\eta} + 2nLT^{\frac{2}{3}} + 4nL^2\eta T^{\frac{4}{3}} + 12LT^{\frac{2}{3}} \sqrt{n^2 + n \log \frac{1}{\delta}} + 6\eta L^2 T \sqrt{n \log \frac{1}{\delta}} \end{aligned}$$

with probability at least  $1 - 3\delta$ . Then, it suffices to bound the term  $\sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T \mathbb{E}[f_t(S^t) \mid x^t]$  using Proposition 3. Consider the random variables  $X_t = f_t(S^t)$  for all  $1 \leq t \leq T$  that are

adapted to the natural filtration generated by the iterates  $\{x^t\}_{t \geq 1}$ . Since  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ , we have  $|X_t| \leq L$ . Further, we have  $\mathbb{E}[X_t^2 | x_t] \leq L^2$ . Applying Proposition 3, we have

$$\mathbb{P} \left( \left| \sum_{t=1}^T f_t(S^t) - \mathbb{E}[f_t(S^t) | x^t] \right| > L \sqrt{2T \log \frac{1}{\delta}} + \frac{L}{2} \log \frac{1}{\delta} \right) \leq \delta$$

Since  $T > \log^{\frac{3}{2}}(1/\delta)$ , we have  $\sqrt{2T \log \frac{1}{\delta}} \geq \frac{1}{2} \log \frac{1}{\delta}$ . This implies that

$$\mathbb{P} \left( \left| \sum_{t=1}^T f_t(S^t) - \mathbb{E}[f_t(S^t) | x^t] \right| > 3L \sqrt{T \log \frac{1}{\delta}} \right) \leq \delta$$

Therefore, we conclude that  $\sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T \mathbb{E}[f_t(S^t) | x^t] \leq 3L \sqrt{T \log \frac{1}{\delta}}$  with probability at least  $1 - \delta$ . Putting these pieces together yields that

$$\begin{aligned} \sum_{t=1}^T f_t(S^t) &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{n}{2\eta} + 3L \sqrt{T \log \frac{1}{\delta}} + 2nLT^{\frac{2}{3}} \\ &\quad + 4nL^2\eta T^{\frac{4}{3}} + 12nLT^{\frac{2}{3}} + 12LT^{\frac{2}{3}} \sqrt{n \log \frac{1}{\delta}} + 6\eta L^2 T \sqrt{n \log \frac{1}{\delta}} \end{aligned}$$

with probability at least  $1 - 4\delta$ . Plugging the choice of  $\eta = \frac{1}{LT^{2/3}}$  yields that

$$\sum_{t=1}^T f_t(S^t) \leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{37}{2} nLT^{\frac{2}{3}} + 21LT^{\frac{2}{3}} \sqrt{n \log \frac{1}{\delta}}$$

with probability at least  $1 - 4\delta$ . Letting  $S = S_\star^T = \operatorname{argmin}_{S \subseteq [n]} \sum_{t=1}^T f_t(S)$  and changing  $\delta$  to  $\frac{\delta}{4}$  yields that  $R_{\alpha,\beta}(T) = O(nT^{\frac{2}{3}} + \sqrt{n \log \frac{1}{\delta}} T^{\frac{2}{3}})$  with probability at least  $1 - \delta$  as desired.  $\square$

## 5 Online Delayed Approximation Algorithm: Full Information Setting

We investigate Algorithm 1 and 2 for regret minimization even when the delay between choosing an action and receiving the incurred cost exists and can be unbounded.

**The general framework.** The general online learning framework with large delay that we consider can be represented as follows. In each round  $t = 1, \dots, T$ , the agent chooses the decision  $S^t \subseteq [n]$  and this generates a loss  $f_t(S^t)$ . Simultaneously,  $S^t$  triggers a *delay*  $d_t \geq 0$  which determines the round  $t + d_t$  at which the information about  $f_t$  will be received. Finally, the agent receives the information about  $f_t$  from all previous rounds  $\mathcal{R}_t = \{s : s + d_s = t\}$ .

The above model has been stated in an abstract way as the basis for the regret analysis. The information about  $f_t$  is determined by whether the setting is full information or bandit feedback. Our blanket assumptions for the stream of the delays encountered will be:

**Assumption 1.** *The delays  $d_t = o(t^\gamma)$  for some  $\gamma < 1$ .*



---

**Algorithm 3** Delay Online Approximate Gradient Descent

---

- 1: **Initialization:** the point  $x^1 \in [0, 1]^n$  and the stepsize  $\eta_t > 0$ ;  $\mathcal{P}_0 \leftarrow \emptyset$  and  $f_\infty = 0$
  - 2: **for**  $t = 1, 2, \dots$  **do**
  - 3:   Let  $x_{\pi(1)}^t \geq \dots x_{\pi(n)}^t$  be the sorted entries in the decreasing order with  $A_i^t = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$  and  $A_0^t = \emptyset$ . Let  $x_{\pi(0)}^t = 1$  and  $x_{\pi(n+1)}^t = 0$
  - 4:   Let  $\lambda_i^t = x_{\pi(i)}^t - x_{\pi(i+1)}^t$  for all  $0 \leq i \leq n$
  - 5:   Sample  $S^t$  from the distribution  $\mathbb{P}(S^t = A_i^t) = \lambda_i^t$  for  $0 \leq i \leq n$  and observe the new loss function  $f_t$
  - 6:   Compute  $g_{\pi(i)}^t = f_t(A_i^t) - f_t(A_{i-1}^t)$  for all  $i \in [n]$  and then trigger a delay  $d_t \geq 0$
  - 7:   Let  $\mathcal{R}_t = \{s : s + d_s = t\}$  and  $\mathcal{P}_t \leftarrow \mathcal{P}_{t-1} \cup \mathcal{R}_t$ . Take  $q_t = \min \mathcal{P}_t$  and set  $\mathcal{P}_t \leftarrow \mathcal{P}_t \setminus \{q_t\}$
  - 8:   Compute  $x^{t+1}$  using Eq. (14)
  - 9: **end for**
- 

Assumption 1 is not theoretically artificial but uncovers that long delays are observed in practice [Chapelle(2014)]; indeed, the data statistics from real-time bidding company suggested that more than 10% of the conversions were  $\geq 2$  weeks old. More specifically, [Chapelle(2014)] showed that the delays in online advertising have long-tail distributions when conditioning on context and feature variables available to the advertiser, thus justifying the existence of unbounded delays. Note that Assumption 1 is mild and the delays can even be adversarial as in [Quanrud and Khashabi(2015)].

**Full information setting.** At the round  $t$ , the agent receives the loss function  $f_s(\cdot)$  for  $\mathcal{R}_t = \{s : s + d_s = t\}$  after committing her decision, i.e., gets to observe  $f_s(A_i^t)$  for all  $s \in \mathcal{R}_t$  and all  $0 \leq i \leq n$ . To let Algorithm 1 handle these delays, the first thing to note is that the set  $\mathcal{R}_t$  received at a given round might be empty, i.e., we could have  $\mathcal{R}_t = \emptyset$  for some  $t \geq 1$ . Following up the pooling strategy in [Héliou et al.(2020)], we assume that, as information is received over time, the agent adds it to an information pool  $\mathcal{P}_t$  and then uses the oldest information available in the pool (where “oldest” stands for the time at which the information was generated).

Since no information is available at  $t = 0$ , we have  $\mathcal{P}_0 = \emptyset$  and update the agent’s information pool recursively:  $\mathcal{P}_t = \mathcal{P}_{t-1} \cup \mathcal{R}_t \setminus \{q_t\}$  where  $q_t = \min(\mathcal{P}_{t-1} \cup \mathcal{R}_t)$  denotes the oldest round from which the agent has unused information at round  $t$ . As [Héliou et al.(2020)] pointed out, this scheme can be seen as a priority queue where  $\{f_s(\cdot), s \in \mathcal{R}_t\}$  arrives at time  $t$  and is assigned in order; subsequently, the oldest information is utilized at first. An important issue that arises in the above computation is that, it may well happen that the agent’s information pool  $\mathcal{P}_t$  is empty at time  $t$  (e.g., if we have  $d_1 > 0$  at time  $t = 1$ ). Following the convention that  $\inf \emptyset = +\infty$ , we set  $q_t = +\infty$  and  $g^\infty = 0$  (since it is impossible to have information at time  $t = +\infty$ ). Under this convention, the computation of a new iterate  $x^{t+1}$  at time  $t$  can be written more explicitly form as follows,

$$x^{t+1} = \begin{cases} x^t & \text{if } \mathcal{P}_t = \emptyset \\ P_{[0,1]^n}(x^t - \eta_t g^{q_t}), & \text{otherwise} \end{cases} \quad (14)$$

We present a delayed variant of Algorithm 1 in Algorithm 3. There is no information aggregation here but the updates of  $x^{t+1}$  follows the pooling policy induced by a priority queue. We summarize our results in the following theorem.

**Theorem 3.** Suppose the adversary chooses nonsubmodular functions in Eq. (8) satisfying  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  and let the delays satisfy Assumption 1. Fixing  $T \geq 1$  and letting  $\eta_t = \frac{\sqrt{n}}{L\sqrt{t^{1+\gamma}}}$  in Algorithm 3, we have  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(\sqrt{nT^{1+\gamma}})$  and  $R_{\alpha,\beta}(T) = O(\sqrt{nT^{1+\gamma}} + \sqrt{T \log \frac{1}{\delta}})$  with probability  $1 - \delta$ .

**Remark.** Theorem 3 demonstrates that Algorithm 3 is no-regret if Assumption 1 hold. To our knowledge, this is the first theoretical guarantee for no-regret learning in online nonsubmodular minimization with delayed costs and also complement similar results for online convex optimization with delayed costs [Quanrud and Khashabi(2015)].

### 5.1 Regret Analysis for Algorithm 3

In this section, we present several technical lemmas for analyzing the regret minimization property of Algorithm 3. We also give the missing proofs of Theorem 3.

We provide one technical lemma for Algorithm 3 which is analogues to Lemma 2. It gives a key inequality for analyzing the regret minimization property of Algorithm 3. Note that the results in Lemma 1 still hold true for the iterates  $\{x^t\}_{t \geq 1}$  and  $\{g^t\}_{t \geq 1}$  generated by Algorithm 3.

**Lemma 6.** Suppose that the iterates  $\{x^t\}_{t \geq 1}$  are generated by Algorithm 3 and  $x \in [0, 1]^n$  and let  $\bar{f}_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ . Then, we have

$$\sum_{t=1}^T \mathbb{E}[(f_t)_L(x^t)] \leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta(-\underline{f}_t)_C(x) \right) + \frac{n}{2\eta_{\bar{T}}} + \frac{L^2}{2} \left( \sum_{t=1}^{\bar{T}} \eta_t \right) + L^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \eta_s \right)$$

where  $\bar{T} > 0$  in the above inequality satisfies that  $q_{\bar{T}} = T$ .

*Proof of Theorem 3.* By [Héliou et al.(2020), Corollary 1], we have  $t - q_t = o(t^\gamma)$  under Assumption 1; in particular, we have  $t - q_t = o(t)$  and  $q_t = \Theta(t)$ . Since  $q_{\bar{T}} = T$ , we have  $T = \Theta(\bar{T})$  which implies that  $\bar{T} = \Theta(T)$ . Recall that  $\eta_t = \frac{\sqrt{n}}{L\sqrt{t^{1+\gamma}}}$ , we have

$$\begin{aligned} \frac{n}{2\eta_{\bar{T}}} &= \frac{L\sqrt{n\bar{T}^{1+\gamma}}}{2} = O(\sqrt{nT^{1+\gamma}}) \\ \frac{L^2}{2} \left( \sum_{t=1}^{\bar{T}} \eta_t \right) &= \frac{\sqrt{n}L}{2} \left( \sum_{t=1}^{\bar{T}} \frac{1}{\sqrt{t^{1+\gamma}}} \right) \leq \frac{\sqrt{n}L}{1-\gamma} \sqrt{\bar{T}^{1-\gamma}} = O(\sqrt{nT^{1-\gamma}}) \\ L^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \eta_s \right) &\leq L^2 \left( \sum_{t=1}^{\bar{T}} (t - q_t) \eta_{q_t} \right) = O \left( \sqrt{n}L \sum_{t=1}^{\bar{T}} \frac{1}{\sqrt{t^{1-\gamma}}} \right) = O(L\sqrt{n\bar{T}^{1+\gamma}}) = O(\sqrt{nT^{1+\gamma}}) \end{aligned}$$

Putting these pieces together with Lemma 6 yields that

$$\sum_{t=1}^T \mathbb{E}[(f_t)_L(x^t)] - \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta(-\underline{f}_t)_C(x) \right) = O(\sqrt{nT^{1+\gamma}}) \quad (15)$$

By the definition of the Lovász extension, we have

$$(f_t)_L(x^t) = \sum_{i=1}^{n-1} (x_{\pi(i)}^t - x_{\pi(i+1)}^t) f_t(A_i^t) + (1 - x_{\pi(1)}^t) f_t(A_0^t) + x_{\pi(n)}^t f_t(A_n^t)$$

---

**Algorithm 4** Delay Bandit Approximate Gradient Descent

---

- 1: **Initialization:** the point  $x^1 \in [0, 1]^n$  and the stepsize  $\eta_t > 0$ ;  $\mathcal{P}_0 \leftarrow \emptyset$  and  $f_\infty = 0$ ; the exploration probability  $\mu_t \in (0, 1)$
  - 2: **for**  $t = 1, 2, \dots$  **do**
  - 3:   Let  $x_{\pi(1)}^t \geq \dots x_{\pi(n)}^t$  be the sorted entries in the decreasing order with  $A_i^t = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$  and  $A_0^t = \emptyset$ . Let  $x_{\pi(0)}^t = 1$  and  $x_{\pi(n+1)}^t = 0$
  - 4:   Let  $\lambda_i^t = x_{\pi(i)}^t - x_{\pi(i+1)}^t$  for all  $0 \leq i \leq n$
  - 5:   Sample  $S^t$  from the distribution  $\mathbb{P}(S^t = A_i^t) = (1 - \mu_t)\lambda_i^t + \frac{\mu_t}{n+1}$  for  $0 \leq i \leq n$  and observe the loss  $f_t(S^t)$
  - 6:   Compute  $\hat{f}_i^t = \frac{\mathbf{1}(S^t = A_i^t)}{(1 - \mu_t)\lambda_i^t + \mu_t/(n+1)} f_t(S^t)$
  - 7:   Compute  $\hat{g}_{\pi(i)}^t = \hat{f}_i^t - \hat{f}_{i-1}^t$  for all  $i \in [n]$  and then trigger a delay  $d_t \geq 0$
  - 8:   Let  $\mathcal{R}_t = \{s : s + d_s = t\}$  and  $\mathcal{P}_t \leftarrow \mathcal{P}_{t-1} \cup \mathcal{R}_t$ . Take  $q_t = \min \mathcal{P}_t$  and set  $\mathcal{P}_t \leftarrow \mathcal{P}_t \setminus \{q_t\}$
  - 9:   Compute  $x^{t+1}$  using Eq. (16)
  - 10: **end for**
- 

By the update formula, we have  $\mathbb{E}[f_t(S^t) \mid x^t] = (f_t)_L(x^t)$  which implies that  $\mathbb{E}[f_t(S^t)] = \mathbb{E}[(f_t)_L(x^t)]$ . Further, by using the same argument as in Theorem 1, we have

$$\frac{1}{\alpha}(\bar{f}_t)_C(\chi_{S_\star^T}) + \beta(-\underline{f}_t)_C(\chi_{S_\star^T}) = \frac{1}{\alpha}\bar{f}_t(S_\star^T) - \beta\underline{f}_t(S_\star^T)$$

Putting these pieces together and letting  $x = \chi_{S_\star^T}$  in Eq. (15) yields that

$$\sum_{t=1}^T \mathbb{E}[f_t(S^t)] - \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T) \right) = O(\sqrt{nT^{1+\gamma}})$$

which implies that  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(\sqrt{nT^{1+\gamma}})$  as desired.

We proceed to derive a high probability bound using the concentration inequality in Proposition 2. Indeed, we have

$$\mathbb{P} \left( \sum_{i=1}^n f_t(S^t) - \mathbb{E} \left[ \sum_{i=1}^n f_t(S^t) \right] > \epsilon \right) \leq \exp \left( -\frac{\epsilon^2}{2nL^2} \right)$$

Equivalently, we have  $\sum_{i=1}^n f_t(S^t) - \mathbb{E}[\sum_{i=1}^n f_t(S^t)] \leq L\sqrt{2T \log \frac{1}{\delta}}$  with probability at least  $1 - \delta$ .

This together with  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(\sqrt{nT^{1+\gamma}})$  yields that  $R_{\alpha,\beta}(T) = O(\sqrt{nT^{1+\gamma}} + \sqrt{T \log \frac{1}{\delta}})$  with probability at least  $1 - \delta$ .  $\square$

## 6 Online Delayed Approximation Algorithm: Bandit Feedback Setting

As we have done in the previous section, we will make use of an unbiased estimator  $\hat{g}$  of the gradient for the bandit feedback setting. However, we only receive the old estimator  $\hat{g}^{q_t}$  at the round  $t$  due

to the delay  $d_t$ . Following the same reasoning as in the full information setting, the computation of a new iterate  $x^{t+1}$  at time  $t$  can be written more explicitly form as follows,

$$x^{t+1} = \begin{cases} x^t & \text{if } \mathcal{P}_t = \emptyset \\ P_{[0,1]^n}(x^t - \eta_t \hat{g}^{q_t}), & \text{otherwise} \end{cases} \quad (16)$$

Algorithm 4 follows the same template as Algorithm 3 but substituting the exact gradients with the gradient estimator. We summarize our results in the following theorem.

**Theorem 4.** *Suppose the adversary chooses nonsubmodular functions in Eq. (8) satisfying  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  and let the delays satisfy Assumption 1. Fixing  $T \geq 1$  and letting  $(\eta_t, \mu_t) = (\frac{1}{Lt^{(2+\gamma)/3}}, \frac{n}{t^{(1-\gamma)/3}})$  in Algorithm 4, we have  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(nT^{\frac{2+\gamma}{3}})$  and  $R_{\alpha,\beta}(T) = O(nT^{\frac{2+\gamma}{3}} + \sqrt{n \log \frac{1}{\delta}} T^{\frac{4-\gamma}{6}})$  with probability  $1 - \delta$ .*

**Remark.** Theorem 4 demonstrates that Algorithm 4 attains the regret of  $nT^{\frac{2+\gamma}{3}}$  which is worse than that of  $\sqrt{n}T^{1+\gamma}$  for Algorithm 3 and reduces to that of  $nT^{\frac{2}{3}}$  for Algorithm 2. Since  $\gamma < 1$  is assumed, Algorithm 4 is the first no-regret bandit learning algorithm for online nonsubmodular minimization with delayed costs to our knowledge.

## 6.1 Regret Analysis for Algorithm 4

In this section, we present several technical lemmas for analyzing the regret minimization property of Algorithm 4. We also give the missing proofs of Theorem 4.

We provide two technical lemmas for Algorithm 4 which are analogues to Lemma 4 and 5. It gives a key inequality for analyzing the regret minimization property of Algorithm 3. Note that the results in Lemma 3 still hold true for the iterates  $\{x^t\}_{t \geq 1}$  and  $\{\hat{g}^t\}_{t \geq 1}$  generated by Algorithm 4.

**Lemma 7.** *Suppose that the iterates  $\{x^t\}_{t \geq 1}$  are generated by Algorithm 4 and  $x \in [0, 1]^n$  and let  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ . Then, we have*

$$\sum_{t=1}^T \mathbb{E}[(f_t)_L(x^t)] \leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) + \frac{n}{2\eta_T} + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right)$$

where  $\bar{T} > 0$  in the above inequality satisfies that  $q_{\bar{T}} = T$ .

Then, we provide our second lemma which significantly generalizes Lemma 7 for deriving the high-probability bounds.

**Lemma 8.** *Suppose that the iterates  $\{x^t\}_{t \geq 1}$  are generated by Algorithm 4 with  $\eta_t = \frac{1}{Lt^{(2+\gamma)/3}}$ ,  $\mu_t = \frac{n}{t^{(1-\gamma)/3}}$  and  $x \in [0, 1]^n$  and let  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ . Fixing a sufficiently small  $\delta \in (0, 1)$  and letting  $T > \log^{\frac{3}{2+\gamma}}(1/\delta)$ . Then, we have*

$$\begin{aligned} \sum_{t=1}^T (f_t)_L(x^t) &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{n}{2\eta_{\bar{T}}} + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \\ &\quad + 12L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n^2 + n \log \frac{1}{\delta}} + 6L \sqrt{n\bar{T} \log \frac{1}{\delta}} \end{aligned}$$

with probability at least  $1 - 3\delta$  where  $\bar{T} > 0$  in the above inequality satisfies that  $q_{\bar{T}} = T$ .

*Proof of Theorem 4.* By [Héliou et al.(2020), Corollary 1], we have  $t - q_t = o(t^\gamma)$  under Assumption 1; in particular, we have  $t - q_t = o(t)$  and  $q_t = \Theta(t)$ . Since  $q_{\bar{T}} = T$ , we have  $T = \Theta(\bar{T})$  which implies that  $\bar{T} = \Theta(T)$ . Recall that  $\eta_t = \frac{1}{Lt^{(2+\gamma)/3}}$  and  $\mu_t = \frac{n}{t^{(1-\gamma)/3}}$ , we have

$$\begin{aligned} \frac{n}{2\eta_{\bar{T}}} &= \frac{nL\bar{T}^{\frac{2+\gamma}{3}}}{2} = O(nT^{\frac{2+\gamma}{3}}) \\ 4n^2L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) &= 4nL \left( \sum_{t=1}^{\bar{T}} \frac{(q_t)^{\frac{1-\gamma}{3}}}{t^{\frac{2+\gamma}{3}}} \right) = O \left( nL \sum_{t=1}^{\bar{T}} \frac{1}{t^{\frac{1+2\gamma}{3}}} \right) = O(nL\bar{T}^{\frac{2-2\gamma}{3}}) = O(nT^{\frac{2-2\gamma}{3}}) \\ 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) &\leq 4L \left( \sum_{t=1}^{\bar{T}} (t - q_t) \frac{\eta_{q_t}}{\mu_t} \right) = O \left( L \sum_{t=1}^{\bar{T}} \frac{1}{t^{\frac{1-\gamma}{3}}} \right) = O(L\bar{T}^{\frac{2+\gamma}{3}}) = O(T^{\frac{2+\gamma}{3}}) \end{aligned}$$

Putting these pieces together with Lemma 7 yields that

$$\sum_{t=1}^T \mathbb{E}[(f_t)_L(x^t)] - \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) = O(nT^{\frac{2+\gamma}{3}}) \quad (17)$$

By using the similar argument as in Theorem 2, we have

$$\mathbb{E}[f_t(S^t) \mid x^t] - (f_t)_L(x^t) \leq L\mu_t \sum_{i=0}^n \left( \frac{1}{n+1} + \lambda_i^t \right) = 2L\mu_t \quad (18)$$

which implies that

$$\sum_{t=1}^T \mathbb{E}[f_t(S^t)] - \mathbb{E}[(f_t)_L(x^t)] \leq 2L \sum_{t=1}^T \mu_t = O(nT^{\frac{2+\gamma}{3}})$$

Using the same argument as in Theorem 1, we have

$$\frac{1}{\alpha} (\bar{f}_t)_C(\chi_{S_\star^T}) + \beta (-\underline{f}_t)_C(\chi_{S_\star^T}) = \frac{1}{\alpha} f_t(S_\star^T) - \beta \underline{f}_t(S_\star^T), \quad \text{where } S_\star^T = \underset{S \subseteq [n]}{\operatorname{argmin}} \sum_{t=1}^T f_t(S)$$

Putting these pieces together and letting  $x = \chi_{S_\star^T}$  in Eq. (17) yields that

$$\sum_{t=1}^T \mathbb{E}[f_t(S^t)] - \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T) \right) = O(nT^{\frac{2+\gamma}{3}})$$

which implies that  $\mathbb{E}[R_{\alpha,\beta}(T)] = O(nT^{\frac{2+\gamma}{3}})$  as desired.

We proceed to derive a high probability bound using Lemma 8. Indeed, we first consider the case of  $T < 2 \log^{\frac{3}{2+\gamma}}(1/\delta)$ . Since  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$ , we have

$$R_{\alpha,\beta}(T) \leq \sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T \left( \frac{1}{\alpha} \bar{f}_t(S_\star^T) - \beta \underline{f}_t(S_\star^T) \right) \leq \left( 1 + \frac{1}{\alpha} + \beta \right) LT = O(T^{\frac{4-\gamma}{6}} \sqrt{\log \frac{1}{\delta}})$$

For the case of  $T \geq 2 \log^{\frac{3}{2+\gamma}}(1/\delta)$ , we obtain by combining Lemma 8 with Eq. (18) that

$$\begin{aligned} \sum_{t=1}^T \mathbb{E}[f_t(S^t) \mid x^t] &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{n}{2\eta_{\bar{T}}} + 2L \left( \sum_{t=1}^T \mu_t \right) + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) \\ &\quad + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) + 12L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n^2 + n \log \frac{1}{\delta}} + 6L\sqrt{n\bar{T} \log \frac{1}{\delta}} \end{aligned}$$

with probability at least  $1 - 3\delta$ . Then, it suffices to bound the term  $\sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T \mathbb{E}[f_t(S^t) \mid x^t]$  using Proposition 3. By using the same argument as in Theorem 2, we have

$$\mathbb{P} \left( \left| \sum_{t=1}^T f_t(S^t) - \mathbb{E}[f_t(S^t) \mid x^t] \right| > 3L\sqrt{T \log \frac{1}{\delta}} \right) \leq \delta$$

which implies that  $\sum_{t=1}^T f_t(S^t) - \sum_{t=1}^T \mathbb{E}[f_t(S^t) \mid x^t] \leq 3L\sqrt{T \log \frac{1}{\delta}}$  with probability at least  $1 - \delta$ . Putting these pieces together yields that

$$\begin{aligned} \sum_{t=1}^T f_t(S^t) &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{n}{2\eta_{\bar{T}}} + 2L \left( \sum_{t=1}^T \mu_t \right) + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) \\ &\quad + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) + 3L\sqrt{T \log \frac{1}{\delta}} + 12L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n^2 + n \log \frac{1}{\delta}} + 6L\sqrt{n\bar{T} \log \frac{1}{\delta}} \end{aligned}$$

with probability at least  $1 - 4\delta$ . Plugging the choices of  $\eta_t = \frac{1}{Lt^{(2+\gamma)/3}}$  and  $\mu_t = \frac{n}{t^{(1-\gamma)/3}}$  and  $\bar{T} = \Theta(T)$  yields that

$$\sum_{t=1}^T f_t(S^t) - \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) = O \left( nT^{\frac{2+\gamma}{3}} + \sqrt{n \log \frac{1}{\delta}} T^{\frac{4-\gamma}{6}} \right)$$

with probability at least  $1 - 4\delta$ . Letting  $S = S_\star^T = \operatorname{argmin}_{S \subseteq [n]} \sum_{t=1}^T f_t(S)$  and changing  $\delta$  to  $\frac{\delta}{4}$  yields that  $R_{\alpha,\beta}(T) = O(nT^{\frac{2+\gamma}{3}} + \sqrt{n \log \frac{1}{\delta}} T^{\frac{4-\gamma}{6}})$  with probability at least  $1 - \delta$  as desired.  $\square$

## 7 Concluding Remarks

This paper studied online nonsubmodular minimization with special structure through the lens of  $(\alpha, \beta)$ -regret and the extension of generic convex relaxation model. We proved that online approximate gradient descent algorithm and its bandit variant adapted for the convex relaxation model could achieve the bounds of  $O(\sqrt{nT})$  and  $O(nT^{\frac{2}{3}})$  in terms of  $(\alpha, \beta)$ -regret respectively. We also investigated the delayed variants of two algorithms and proved new regret bounds where the delays can even be unbounded. More specifically, if delays satisfy  $d_t = o(t^\gamma)$  with  $\gamma < 1$ , we showed that our proposed algorithms achieve the regret bound of  $O(\sqrt{nT^{1+\gamma}})$  and  $O(nT^{\frac{2+\gamma}{3}})$  for full-information setting and bandit setting respectively. Simulation studies validate our theoretical findings in practice.

## References

- [Anari et al.(2019)] N. Anari, N. Haghtalab, S. Naor, S. Pokutta, M. Singh, and A. Torrico. Structured robust submodular maximization: Offline and online algorithms. In *AISTATS*, pages 3128–3137. PMLR, 2019.
- [Arora et al.(2012)] S. Arora, E. Hazan, and S. Kale. The multiplicative weights update method: A meta-algorithm and applications. *Theory of Computing*, 8(1):121–164, 2012.
- [Auer(2002)] P. Auer. Using confidence bounds for exploitation-exploration trade-offs. *Journal of Machine Learning Research*, 3(Nov):397–422, 2002.
- [Bach(2010)] F. Bach. Structured sparsity-inducing norms through submodular functions. In *NeurIPS*, pages 118–126, 2010.
- [Bai et al.(2016)] W. Bai, R. Iyer, K. Wei, and J. Bilmes. Algorithms for optimizing the ratio of submodular functions. In *ICML*, pages 2751–2759. PMLR, 2016.
- [Bian et al.(2017)] A. A. Bian, J. M. Buhmann, A. Krause, and S. Tschieschek. Guarantees for greedy maximization of non-submodular functions with applications. In *ICML*, pages 498–507. PMLR, 2017.
- [Bistritz et al.(2019)] I. Bistritz, Z. Zhou, X. Chen, N. Bambos, and J. Blanchet. Online EXP3 learning in adversarial bandits with delayed feedback. In *NeurIPS*, page 11349–11358, 2019.
- [Blum and Mansour(2007)] A. Blum and Y. Mansour. From external to internal regret. *Journal of Machine Learning Research*, 8(6), 2007.
- [Bogunovic et al.(2016)] I. Bogunovic, J. Scarlett, A. Krause, and V. Cevher. Truncated variance reduction: A unified approach to Bayesian optimization and level-set estimation. In *NeurIPS*, pages 1507–1515, 2016.
- [Bogunovic et al.(2018)] I. Bogunovic, J. Zhao, and V. Cevher. Robust maximization of non-submodular objectives. In *AISTATS*, pages 890–899. PMLR, 2018.
- [Boykov and Kolmogorov(2004)] Y. Boykov and V. Kolmogorov. An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(9):1124–1137, 2004.
- [Boykov et al.(2001)] Y. Boykov, O. Veksler, and R. Zabih. Fast approximate energy minimization via graph cuts. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 23(11):1222–1239, 2001.
- [Buchbinder et al.(2014)] N. Buchbinder, M. Feldman, and R. Schwartz. Online submodular maximization with preemption. In *SODA*, pages 1202–1216. SIAM, 2014.
- [Cardoso and Cummings(2019)] A. R. Cardoso and R. Cummings. Differentially private online submodular minimization. In *AISTATS*, pages 1650–1658. PMLR, 2019.
- [Cesa-Bianchi and Lugosi(2006)] N. Cesa-Bianchi and G. Lugosi. *Prediction, Learning, and Games*. Cambridge University Press, 2006.
- [Cesa-Bianchi and Lugosi(2012)] N. Cesa-Bianchi and G. Lugosi. Combinatorial bandits. *Journal of Computer and System Sciences*, 78(5):1404–1422, 2012.
- [Chapelle(2014)] O. Chapelle. Modeling delayed feedback in display advertising. In *KDD*, pages 1097–1105, 2014.
- [Chen et al.(2018a)] L. Chen, M. Feldman, and A. Karbasi. Weakly submodular maximization beyond cardinality constraints: Does randomization help greedy? In *ICML*, pages 804–813. PMLR, 2018a.
- [Chen et al.(2018b)] L. Chen, C. Harshaw, H. Hassani, and A. Karbasi. Projection-free online optimization with stochastic gradient: From convexity to submodularity. In *ICML*, pages 814–823. PMLR, 2018b.

- [Chen et al.(2018c)] L. Chen, H. Hassani, and A. Karbasi. Online continuous submodular maximization. In *AISTATS*, pages 1896–1905. PMLR, 2018c.
- [Conforti and Cornuéjols(1984)] M. Conforti and G. Cornuéjols. Submodular set functions, matroids and the greedy algorithm: tight worst-case bounds and some generalizations of the Rado-Edmonds theorem. *Discrete Applied Mathematics*, 7(3):251–274, 1984.
- [Dani et al.(2008)] V. Dani, T. P. Hayes, and S. M. Kakade. Stochastic linear optimization under bandit feedback. In *COLT*, pages 355–366. Omnipress, 2008.
- [Das and Kempe(2011)] A. Das and D. Kempe. Submodular meets spectral: Greedy algorithms for subset selection, sparse approximation and dictionary selection. In *ICML*, pages 1057–1064, 2011.
- [Dughmi(2009)] S. Dughmi. Submodular functions: Extensions, distributions, and algorithms. a survey. *ArXiv Preprint: 0912.0322*, 2009.
- [Edmonds(2003)] J. Edmonds. Submodular functions, matroids, and certain polyhedra. In *Combinatorial Optimization—Eureka, You Shrink!*, pages 11–26. Springer, 2003.
- [El Halabi and Cevher(2015)] M. El Halabi and V. Cevher. A totally unimodular view of structured sparsity. In *AISTATS*, pages 223–231. PMLR, 2015.
- [El Halabi and Jegelka(2020)] M. El Halabi and S. Jegelka. Optimal approximation for unconstrained non-submodular minimization. In *ICML*, pages 3961–3972. PMLR, 2020.
- [El Halabi et al.(2018)] M. El Halabi, F. Bach, and V. Cevher. Combinatorial penalties: Which structures are preserved by convex relaxations? In *AISTATS*, pages 1551–1560. PMLR, 2018.
- [Elenberg et al.(2018)] E. R. Elenberg, R. Khanna, A. G. Dimakis, and S. Negahban. Restricted strong convexity implies weak submodularity. *The Annals of Statistics*, 46(6B):3539–3568, 2018.
- [Flaxman et al.(2005)] A. D. Flaxman, A. T. Kalai, and H. B. McMahan. Online convex optimization in the bandit setting: gradient descent without a gradient. In *SODA*, pages 385–394, 2005.
- [Freedman(1975)] D. A. Freedman. On tail probabilities for martingales. *The Annals of Probability*, pages 100–118, 1975.
- [Fujishige(2005)] S. Fujishige. *Submodular Functions and Optimization*. Elsevier, 2005.
- [González et al.(2016)] J. González, Z. Dai, P. Hennig, and N. Lawrence. Batch Bayesian optimization via local penalization. In *AISTATS*, pages 648–657. PMLR, 2016.
- [Grötschel et al.(2012)] M. Grötschel, L. Lovász, and A. Schrijver. *Geometric algorithms and combinatorial optimization*, volume 2. Springer Science & Business Media, 2012.
- [Gyorgy and Joulani(2021)] A. Gyorgy and P. Joulani. Adapting to delays and data in adversarial multi-armed bandits. In *ICML*, pages 3988–3997. PMLR, 2021.
- [Harshaw et al.(2019)] C. Harshaw, M. Feldman, J. Ward, and A. Karbasi. Submodular maximization beyond non-negativity: Guarantees, fast algorithms, and applications. In *ICML*, pages 2634–2643. PMLR, 2019.
- [Harvey et al.(2020)] N. Harvey, C. Liaw, and T. Soma. Improved algorithms for online submodular maximization via first-order regret bounds. In *NeurIPS*, pages 123–133. Curran Associates, Inc., 2020.
- [Hassidim and Singer(2018)] A. Hassidim and Y. Singer. Optimization for approximate submodularity. In *NeurIPS*, pages 394–405, 2018.
- [Hazan(2016)] E. Hazan. Introduction to online convex optimization. *Foundations and Trends in Optimization*, 2(3-4):157–325, 2016.
- [Hazan and Kale(2012)] E. Hazan and S. Kale. Online submodular minimization. *Journal of Machine Learning Research*, 13(10), 2012.



- [Hazan et al.(2007)] E. Hazan, A. Agarwal, and S. Kale. Logarithmic regret algorithms for online convex optimization. *Machine Learning*, 69(2-3):169–192, 2007.
- [Héliou et al.(2020)] A. Héliou, P. Mertikopoulos, and Z. Zhou. Gradient-free online learning in continuous games with delayed rewards. In *ICML*, pages 4172–4181. PMLR, 2020.
- [Hoeffding(1963)] W. Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.
- [Horel and Singer(2016)] T. Horel and Y. Singer. Maximization of approximately submodular functions. In *NeurIPS*, pages 3045–3053, 2016.
- [Iwata(2003)] S. Iwata. A faster scaling algorithm for minimizing submodular functions. *SIAM Journal on Computing*, 32(4):833–840, 2003.
- [Iwata and Orlin(2009)] S. Iwata and J. B. Orlin. A simple combinatorial algorithm for submodular function minimization. In *SODA*, pages 1230–1237, 2009.
- [Iwata et al.(2001)] S. Iwata, L. Fleischer, and S. Fujishige. A combinatorial strongly polynomial algorithm for minimizing submodular functions. *Journal of the ACM*, 48(4):761–777, 2001.
- [Iyer and Bilmes(2012)] R. Iyer and J. Bilmes. Algorithms for approximate minimization of the difference between submodular functions, with applications. In *UAI*, pages 407–417, 2012.
- [Iyer et al.(2013)] R. K. Iyer, S. Jegelka, and J. A. Bilmes. Curvature and optimal algorithms for learning and minimizing submodular functions. In *NeurIPS*, pages 2742–2750, 2013.
- [Jegelka and Bilmes(2011)] S. Jegelka and J. Bilmes. Online submodular minimization for combinatorial structures. In *ICML*, pages 345–352, 2011.
- [Joulani et al.(2013)] P. Joulani, A. Gyorgy, and C. Szepesvári. Online learning under delayed feedback. In *ICML*, pages 1453–1461. PMLR, 2013.
- [Joulani et al.(2016)] P. Joulani, A. Gyorgy, and C. Szepesvári. Delay-tolerant online convex optimization: Unified analysis and adaptive-gradient algorithms. In *AAAI*, pages 1744–1750, 2016.
- [Kalai and Vempala(2005)] A. Kalai and S. Vempala. Efficient algorithms for online decision problems. *Journal of Computer and System Sciences*, 71(3):291–307, 2005.
- [Kawahara et al.(2015)] Y. Kawahara, R. Iyer, and J. Bilmes. On approximate non-submodular minimization via tree-structured supermodularity. In *AISTATS*, pages 444–452. PMLR, 2015.
- [Kuhnle et al.(2018)] A. Kuhnle, J. D. Smith, V. Crawford, and M. Thai. Fast maximization of non-submodular, monotonic functions on the integer lattice. In *ICML*, pages 2786–2795. PMLR, 2018.
- [Lattimore and Szepesvári(2020)] T. Lattimore and C. Szepesvári. *Bandit Algorithms*. Cambridge University Press, 2020.
- [Lee et al.(2015)] Y. T. Lee, A. Sidford, and S. C-W. Wong. A faster cutting plane method and its implications for combinatorial and convex optimization. In *FOCS*, pages 1049–1065. IEEE, 2015.
- [Lehmann et al.(2006)] B. Lehmann, D. Lehmann, and N. Nisan. Combinatorial auctions with decreasing marginal utilities. *Games and Economic Behavior*, 55(2):270–296, 2006.
- [Lovász(1983)] L. Lovász. Submodular functions and convexity. In *Mathematical Programming The State of The Art*, pages 235–257. Springer, 1983.
- [Matsuoka et al.(2021)] T. Matsuoka, S. Ito, and N. Ohsaka. Tracking regret bounds for online submodular optimization. In *AISTATS*, pages 3421–3429. PMLR, 2021.
- [McCormick(2005)] S. T. McCormick. Submodular function minimization. *Handbooks in Operations Research and Management Science*, 12:321–391, 2005.

- [Narasimhan et al.(2005)] M. Narasimhan, N. Jojic, and J. Bilmes. Q-Clustering. In *NeurIPS*, pages 979–986, 2005.
- [Orlin(2009)] J. B. Orlin. A faster strongly polynomial time algorithm for submodular function minimization. *Mathematical Programming*, 118(2):237–251, 2009.
- [Pike-Burke et al.(2018)] C. Pike-Burke, S. Agrawal, C. Szepesvari, and S. Grunewalder. Bandits with delayed, aggregated anonymous feedback. In *ICML*, pages 4105–4113. PMLR, 2018.
- [Quanrud and Khashabi(2015)] K. Quanrud and D. Khashabi. Online learning with adversarial delays. In *NeurIPS*, pages 1270–1278, 2015.
- [Rapaport et al.(2008)] F. Rapaport, E. Barillot, and J-P. Vert. Classification of arrayCGH data using fused SVM. *Bioinformatics*, 24(13):i375–i382, 2008.
- [Roughgarden and Wang(2018)] T. Roughgarden and J. R. Wang. An optimal learning algorithm for online unconstrained submodular maximization. In *COLT*, pages 1307–1325. PMLR, 2018.
- [Sakaue(2019)] S. Sakaue. Greedy and IHT algorithms for non-convex optimization with monotone costs of non-zeros. In *AISTATS*, pages 206–215. PMLR, 2019.
- [Schrijver(2000)] A. Schrijver. A combinatorial algorithm minimizing submodular functions in strongly polynomial time. *Journal of Combinatorial Theory, Series B*, 80(2):346–355, 2000.
- [Shalev-Shwartz(2011)] S. Shalev-Shwartz. Online learning and online convex optimization. *Foundations and Trends in Machine Learning*, 4(2):107–194, 2011.
- [Shalev-Shwartz and Singer(2006)] S. Shalev-Shwartz and Y. Singer. Convex repeated games and Fenchel duality. In *NIPS*, pages 1265–1272, 2006.
- [Streeter and Golovin(2008)] M. Streeter and D. Golovin. An online algorithm for maximizing submodular functions. In *NeurIPS*, pages 1577–1584, 2008.
- [Sviridenko et al.(2017)] M. Sviridenko, J. Vondrák, and J. Ward. Optimal approximation for submodular and supermodular optimization with bounded curvature. *Mathematics of Operations Research*, 42(4):1197–1218, 2017.
- [Svitkina and Fleischer(2011)] Z. Svitkina and L. Fleischer. Submodular approximation: Sampling-based algorithms and lower bounds. *SIAM Journal on Computing*, 40(6):1715–1737, 2011.
- [Thang and Srivastav(2021)] N. K. Thang and A. Srivastav. Online non-monotone DR-submodular maximization. In *AAAI*, pages 9868–9876, 2021.
- [Thune et al.(2019)] T. S. Thune, N. Cesa-Bianchi, and Y. Seldin. Nonstochastic multiarmed bandits with unrestricted delays. In *NeurIPS*, pages 6538–6547, 2019.
- [Trevisan(2014)] L. Trevisan. Inapproximability of combinatorial optimization problems. *Paradigms of Combinatorial Optimization: Problems and New Approaches*, pages 381–434, 2014.
- [Vernade et al.(2017)] C. Vernade, O. Cappé, and V. Perchet. Stochastic bandit models for delayed conversions. In *UAI*. AUAI Press, 2017.
- [Vondrák(2007)] J. Vondrák. *Submodularity in Combinatorial Optimization*. PhD thesis, Charles University, 2007.
- [Zhou et al.(2017)] Z. Zhou, P. Mertikopoulos, N. Bambos, P. Glynn, and C. Tomlin. Countering feedback delays in multi-agent learning. In *NeurIPS*, pages 6172–6182, 2017.
- [Zhou et al.(2019)] Z. Zhou, R. Xu, and J. Blanchet. Learning in generalized linear contextual bandits with stochastic delays. In *NeurIPS*, pages 5197–5208, 2019.
- [Zimmert and Seldin(2020)] J. Zimmert and Y. Seldin. An optimal algorithm for adversarial bandits with arbitrary delays. In *AISTATS*, pages 3285–3294. PMLR, 2020.
- [Zinkevich(2003)] M. Zinkevich. Online convex programming and generalized infinitesimal gradient ascent. In *ICML*, pages 928–936, 2003.

## A Proof of Auxiliary Lemmas

*Proof of Lemma 1.* Since  $x^t \in [0, 1]^n$  and  $x \in [0, 1]^n$  is fixed, we have  $\|x^t - x\| \leq \sqrt{\sum_{i=1}^n 1} = \sqrt{n}$  for all  $t \geq 1$ . By the definition of  $g^t$ , we have  $g_{\pi(i)}^t = f_t(A_i^t) - f_t(A_{i-1}^t)$  for all  $i \in [n]$  where  $A_i^t = \{\pi(1), \dots, \pi(i)\}$  for all  $i \in [n]$ . Then, we have

$$\|g^t\| \leq \sum_{i=1}^n |f_t(A_i^t) - f_t(A_{i-1}^t)| \leq \sum_{i=1}^n |\bar{f}_t(A_i^t) - \bar{f}_t(A_{i-1}^t)| + \sum_{i=1}^n |\underline{f}_t(A_i^t) - \underline{f}_t(A_{i-1}^t)|$$

Since  $\bar{f}_t$  and  $\underline{f}_t$  are both normalized and non-decreasing, we have

$$\sum_{i=1}^n |\bar{f}_t(A_i^t) - \bar{f}_t(A_{i-1}^t)| + \sum_{i=1}^n |\underline{f}_t(A_i^t) - \underline{f}_t(A_{i-1}^t)| = \bar{f}_t([n]) + \underline{f}_t([n]) \leq L$$

Putting these pieces together yields that  $\|g^t\| \leq L$  for all  $t = 1, 2, \dots, T$ .  $\square$

*Proof of Lemma 2.* Since  $x^{t+1} = P_{[0,1]^n}(x^t - \eta g^t)$ , we have

$$(x - x^{t+1})^\top (x^{t+1} - x^t + \eta g^t) \geq 0, \quad \text{for all } x \in [0, 1]^n$$

Rearranging the above inequality and using the fact that  $\eta > 0$ , we have

$$(x^{t+1} - x)^\top g^t \leq \frac{1}{\eta} (x - x^{t+1})^\top (x^{t+1} - x^t) = \frac{1}{2\eta} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2 - \|x^t - x^{t+1}\|^2) \quad (19)$$

Using Young's inequality, we have

$$(x^t - x^{t+1})^\top g^t \leq \frac{1}{2\eta} \|x^t - x^{t+1}\|^2 + \frac{\eta}{2} \|g^t\|^2 \quad (20)$$

Combining Eq. (19) and Eq. (20) yields that

$$(x^t - x)^\top g^t \leq \frac{1}{2\eta} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta}{2} \|g^t\|^2$$

Since  $f_t = \bar{f}_t - \underline{f}_t$  where  $\bar{f}_t$  and  $\underline{f}_t$  are both non-decreasing,  $\bar{f}_t$  is  $\alpha$ -weakly DR-submodular and  $\underline{f}_t$  is  $\beta$ -weakly DR-supermodular, Proposition 1 implies that

$$(x^t - x)^\top g^t \geq (f_t)_L(x^t) - \left( \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right)$$

By Lemma 1, we have  $\|g^t\| \leq L$  for all  $t = 1, 2, \dots, T$ . Then, we have

$$(f_t)_L(x^t) \leq \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) + \frac{1}{2\eta} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta L^2}{2}$$

Summing up the above inequality over  $t = 1, 2, \dots, T$  and using  $\|x^1 - x\| \leq \sqrt{n}$  (cf. Lemma 1), we have

$$\sum_{t=1}^T (f_t)_L(x^t) \leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) + \frac{n}{2\eta} + \frac{\eta L^2 T}{2}$$

Taking the expectation of both sides yields the desired inequality.  $\square$

*Proof of Lemma 3.* Using the same argument as in Lemma 1, we have  $\|x^t - x\| \leq \sqrt{n}$  for all  $t \geq 1$ . By the definition of  $\hat{g}^t$ , we have

$$\hat{g}_{\pi(i)}^t = \left( \frac{\mathbf{1}(S^t = A_i^t)}{(1-\mu)\lambda_i^t + \frac{\mu}{n+1}} - \frac{\mathbf{1}(S^t = A_{i-1}^t)}{(1-\mu)\lambda_{i-1}^t + \frac{\mu}{n+1}} \right) f_t(S^t), \quad \text{for all } i \in [n]$$

This together with the sampling scheme for  $S^t$  implies that

$$\mathbb{E}[\hat{g}_{\pi(i)}^t \mid x^t] = f_t(A_i^t) - f_t(A_{i-1}^t), \quad \text{for all } i \in [n]$$

Since  $g_{\pi(i)}^t = f_t(A_i^t) - f_t(A_{i-1}^t)$  for all  $i \in [n]$ , we have  $\mathbb{E}[\hat{g}^t \mid x^t] = g^t$ . Since  $f_t = \bar{f}_t - \underline{f}_t$  satisfy that  $\bar{f}_t([n]) + \underline{f}_t([n]) \leq L$  for all  $t \geq 1$  and  $\bar{f}_t$  and  $\underline{f}_t$  are both normalized and non-decreasing, we have

$$\mathbb{E}[\|\hat{g}^t\|^2 \mid x^t] \leq \sum_{i=0}^n \frac{2(f_t(A_i^t))^2}{(1-\mu)\lambda_i^t + \frac{\mu}{n+1}} \leq \frac{2(n+1)^2 L^2}{\mu} \leq \frac{8n^2 L^2}{\mu}$$

Further, let  $S^t = A_{i_t}$  in the round  $t$ , we can apply the same argument and obtain that

$$\|\hat{g}^t\|^2 \leq 2 \left( \frac{f_t(A_{i_t}^t)}{(1-\mu)\lambda_{i_t}^t + \frac{\mu}{n+1}} \right)^2 \leq \frac{2(n+1)^2 L^2}{\mu^2}$$

This completes the proof.  $\square$

*Proof of Lemma 4.* Using the same argument as in Lemma 2, we have

$$(x^t - x)^\top \hat{g}^t \leq \frac{1}{2\eta} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta}{2} \|\hat{g}^t\|^2$$

By Lemma 3, we have  $\mathbb{E}[\hat{g}^t \mid x^t] = g^t$  and  $\mathbb{E}[\|\hat{g}^t\|^2 \mid x^t] \leq \frac{8n^2 L^2}{\mu}$  for all  $t \geq 1$ . This implies that

$$(x^t - x)^\top g^t \leq \frac{1}{2\eta} (\|x - x^t\|^2 - \mathbb{E}[\|x - x^{t+1}\|^2 \mid x^t]) + \frac{4n^2 L^2 \eta}{\mu}$$

Since  $f_t = \bar{f}_t - \underline{f}_t$  where  $\bar{f}_t$  and  $\underline{f}_t$  are both non-decreasing,  $\bar{f}_t$  is  $\alpha$ -weakly DR-submodular and  $\underline{f}_t$  is  $\beta$ -weakly DR-supermodular, Proposition 1 implies that

$$(x^t - x)^\top g^t \geq (f_t)_L(x^t) - \left( \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right)$$

By Lemma 1, we have  $\|g^t\| \leq L$  for all  $t = 1, 2, \dots, T$ . Then, we have

$$(f_t)_L(x^t) \leq \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) + \frac{1}{2\eta} (\|x - x^t\|^2 - \mathbb{E}[\|x - x^{t+1}\|^2 \mid x^t]) + \frac{4n^2 L^2 \eta}{\mu}$$

Taking the expectation of both sides and summing up the resulting inequality over  $t = 1, 2, \dots, T$ , we have

$$\sum_{t=1}^T \mathbb{E}[(f_t)_L(x^t)] \leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) + \frac{1}{2\eta} \|x - x^1\|^2 + \frac{4n^2 L^2 \eta T}{\mu}$$

Using  $\|x^1 - x\| \leq \sqrt{n}$  (cf. Lemma 3) yields the desired inequality.  $\square$

*Proof of Lemma 5.* Using the same argument as in Lemma 4, we have

$$(x^t - x)^\top \hat{g}^t \leq \frac{1}{2\eta} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta}{2} \|\hat{g}^t\|^2$$

and

$$(x^t - x)^\top g^t \geq (f_t)_L(x^t) - \left( \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right)$$

For simplicity, we define  $e_t = \hat{g}^t - g^t$ . Then, we have

$$(f_t)_L(x^t) - \left( \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) \leq (x - x^t)^\top e_t + \frac{1}{2\eta} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta}{2} \|\hat{g}^t\|^2 \quad (21)$$

Summing up Eq. (21) over  $t = 1, 2, \dots, T$  and using  $\|x^1 - x\| \leq \sqrt{n}$  and  $\mathbb{E}[\|\hat{g}^t\|^2 \mid x^t] \leq \frac{8n^2 L^2}{\mu}$  for all  $t \geq 1$  (cf. Lemma 3), we have

$$\begin{aligned} \sum_{t=1}^T (f_t)_L(x^t) &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) + (x - x^t)^\top e_t + \frac{\eta}{2} (\|\hat{g}^t\|^2 - \mathbb{E}[\|\hat{g}^t\|^2 \mid x^t]) \right) \\ &\quad + \frac{n}{2\eta} + \frac{4n^2 L^2 \eta T}{\mu} \end{aligned}$$

By the definition of the convex closure, we obtain that the convex closure of a set function  $f$  agrees with  $f$  on all the integer points [Dughmi(2009), Page 4, Proposition 3.3]. Letting  $S \subseteq [n]$ , we have  $(\bar{f}_t)_C(\chi_S) = f_t(S)$  and  $(-\underline{f}_t)_C(\chi_S) = -\beta f_t(S)$  which implies that

$$\frac{1}{\alpha} (\bar{f}_t)_C(\chi_S) + \beta (-\underline{f}_t)_C(\chi_S) = \frac{1}{\alpha} f_t(S) - \beta f_t(S)$$

Letting  $x = \chi_S$ , we have

$$\sum_{t=1}^T (f_t)_L(x^t) \leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta f_t(S) \right) + \frac{n}{2\eta} + \frac{4n^2 L^2 \eta T}{\mu} + \underbrace{\sum_{t=1}^T (\chi_S - x^t)^\top e_t}_{\mathbf{I}} + \underbrace{\frac{\eta}{2} \left( \sum_{t=1}^T \|\hat{g}^t\|^2 - \mathbb{E}[\|\hat{g}^t\|^2 \mid x^t] \right)}_{\mathbf{II}} \quad (22)$$

In what follows, we prove the high probability bounds for the terms **I** and **II** in the above inequality.

**Bounding I.** Consider the random variables  $X_t = (x^t)^\top \hat{g}^t$  for all  $1 \leq t \leq T$  that are adapted to the natural filtration generated by the iterates  $\{x_t\}_{t \geq 1}$ . By Lemma 3 and the Hölder's inequality, we have

$$|X_t| \leq \|\hat{g}^t\|_1 \|x^t\|_\infty \leq 2 \left| \frac{f_t(A_{i_t}^t)}{(1 - \mu)\lambda_{i_t}^t + \frac{\mu}{n+1}} \right| \leq \frac{2(n+1)L}{\mu}$$

Since  $\mu = \frac{n}{T^{1/3}}$ , we have  $|X_t| \leq 4LT^{\frac{1}{3}}$ . Further, we have

$$\mathbb{E}[X_t^2 \mid x_t] \leq \mathbb{E}[\|\hat{g}^t\|_1^2 \|x^t\|_\infty^2 \mid x_t] = \sum_{i=0}^n \frac{4(f_t(A_i^t))^2}{(1 - \mu)\lambda_i^t + \frac{\mu}{n+1}} \leq \frac{2(n+1)^2 L^2}{\mu} \leq 8nL^2 T^{\frac{1}{3}}$$

Since  $\mathbb{E}[\hat{g}^t | x^t] = g^t$  and  $e_t = \hat{g}^t - g^t$ , Proposition 3 implies that

$$\mathbb{P} \left( \left| \sum_{t=1}^T (x^t)^\top e_t \right| > 4LT^{\frac{2}{3}} \sqrt{n \log \frac{1}{\delta}} + 2LT^{\frac{1}{3}} \log \frac{1}{\delta} \right) \leq \delta$$

Since  $T > \log^{\frac{3}{2}}(1/\delta)$ , we have  $T^{\frac{2}{3}} \sqrt{\log \frac{1}{\delta}} \geq T^{\frac{1}{3}} \log \frac{1}{\delta}$ . This implies that

$$\mathbb{P} \left( \left| \sum_{t=1}^T (x^t)^\top e_t \right| > 6LT^{\frac{2}{3}} \sqrt{n \log \frac{1}{\delta}} \right) \leq \delta$$

Similarly, we fix a set  $S \subseteq [n]$  and consider the random variable  $X_t = (\chi_S)^\top \hat{g}^t$  for all  $1 \leq t \leq T$  that are adapted to the natural filtration generated by the iterates  $\{x_t\}_{t \geq 1}$ . By repeating the above argument with  $\frac{\delta}{2^n}$ , we have

$$\mathbb{P} \left( \left| \sum_{t=1}^T (\chi_S)^\top e_t \right| > 6LT^{\frac{2}{3}} \sqrt{n \log(2^n/\delta)} \right) \leq \frac{\delta}{2^n}$$

By taking a union bound over the  $2^n$  choices of  $S$ , we obtain that

$$\mathbb{P} \left( \left| \sum_{t=1}^T (\chi_S)^\top e_t \right| > 6LT^{\frac{2}{3}} \sqrt{n \log(2^n/\delta)} \right) \leq \delta, \quad \text{for any } S \subseteq [n]$$

Since  $\sqrt{n \log(2^n/\delta)} \leq \sqrt{n^2 + n \log \frac{1}{\delta}}$ , we have  $\mathbf{I} \leq 12LT^{\frac{2}{3}} \sqrt{n^2 + n \log \frac{1}{\delta}}$  with probability at least  $1 - 2\delta$ .

**Bounding II.** Consider the random variables  $X_t = \|\hat{g}^t\|^2$  for all  $1 \leq t \leq T$  that are adapted to the natural filtration generated by the iterates  $\{x_t\}_{t \geq 1}$ . By Lemma 3, we have  $|X_t| \leq \frac{2(n+1)^2 L^2}{\mu^2}$ . Since  $\mu = \frac{n}{T^{1/3}}$ , we have  $|X_t| \leq 8L^2 T^{2/3}$ . Further, we have

$$\mathbb{E}[X_t^2 | x_t] \leq \sum_{i=0}^n \frac{2(f_t(A_i^t))^4}{((1-\mu)\lambda_i^t + \frac{\mu}{n+1})^3} \leq \frac{2(n+1)^4 L^4}{\mu^3} \leq 32nL^4 T$$

Applying Proposition 3, we have

$$\mathbb{P} \left( \left| \sum_{t=1}^T \|\hat{g}^t\|^2 - \mathbb{E}[\|\hat{g}^t\|^2 | x^t] \right| > 8L^2 T \sqrt{n \log \frac{1}{\delta}} + 4L^2 T^{\frac{2}{3}} \log \frac{1}{\delta} \right) \leq \delta$$

Since  $T > \log^{\frac{3}{2}}(1/\delta)$ , we have  $T \sqrt{\log \frac{1}{\delta}} \geq T^{\frac{2}{3}} \log \frac{1}{\delta}$ . This implies that

$$\mathbb{P} \left( \left| \sum_{t=1}^T \|\hat{g}^t\|^2 - \mathbb{E}[\|\hat{g}^t\|^2 | x^t] \right| > 12L^2 T \sqrt{n \log \frac{1}{\delta}} \right) \leq \delta$$

Therefore, we conclude that  $\mathbf{II} \leq 12L^2 T \sqrt{n \log \frac{1}{\delta}}$  with probability at least  $1 - \delta$ .

Putting these pieces together with Eq. (22) yields that

$$\sum_{t=1}^T (f_t)_L(x^t) \leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{n}{2\eta} + \frac{4n^2 L^2 \eta T}{\mu} + 12LT^{\frac{2}{3}} \sqrt{n^2 + n \log \frac{1}{\delta}} + 6\eta L^2 T \sqrt{n \log \frac{1}{\delta}}$$

with probability at least  $1 - 3\delta$ .  $\square$

*Proof of Lemma 6.* Using the same argument as in Lemma 2, we have

$$(x^t - x)^\top g^{q_t} \leq \frac{1}{2\eta_t} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta_t}{2} \|g^{q_t}\|^2$$

Since  $f_t = \bar{f}_t - \underline{f}_t$  where  $\bar{f}_t$  and  $\underline{f}_t$  are both normalized and non-decreasing,  $\bar{f}_t$  is  $\alpha$ -weakly DR-submodular and  $\underline{f}_t$  is  $\beta$ -weakly DR-supermodular, Proposition 1 implies that

$$(x^{q_t} - x)^\top g^{q_t} \geq (f_{q_t})_L(x^{q_t}) - \left( \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) \right)$$

By Lemma 1, we have  $\|g^t\| \leq L$  for all  $t \geq 1$ . Then, we have

$$(f_{q_t})_L(x^{q_t}) \leq \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) + \frac{1}{2\eta_t} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + L\|x^{q_t} - x^t\| + \frac{\eta_t L^2}{2} \quad (23)$$

Further, we have

$$\|x^{q_t} - x^t\| \leq \sum_{s=q_t}^{t-1} \eta_s \|g^s\| \leq L \left( \sum_{s=q_t}^{t-1} \eta_s \right) \quad (24)$$

Plugging Eq. (24) into Eq. (23) yields that

$$(f_{q_t})_L(x^{q_t}) \leq \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) + \frac{1}{2\eta_t} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta_t L^2}{2} + L^2 \left( \sum_{s=q_t}^{t-1} \eta_s \right) \quad (25)$$

For a fixed horizon  $T \geq 1$ , we have  $q_{\bar{T}} = T$  for some  $\bar{T} \geq T$ . Then, by summing up Eq. (25) over  $t = 1, 2, \dots, \bar{T}$  and using  $\|x^t - x\| \leq \sqrt{n}$  for all  $t \geq 1$  (cf. Lemma 1) and that  $\{\eta_t\}_{t \geq 1}$  is nonincreasing, we have

$$\sum_{t=1}^{\bar{T}} (f_{q_t})_L(x^{q_t}) \leq \left( \sum_{t=1}^{\bar{T}} \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) \right) + \frac{n}{2\eta_{\bar{T}}} + \frac{L^2}{2} \left( \sum_{t=1}^{\bar{T}} \eta_t \right) + L^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \eta_s \right)$$

Since  $q_{\bar{T}} = T$  and our pooling policy is induced by a priority queue (note that  $f_{q_t} = \bar{f}_{q_t} = \underline{f}_{q_t} = 0$  if  $\mathcal{P}_t = \emptyset$ ), we have

$$\begin{aligned} \sum_{t=1}^{\bar{T}} (f_{q_t})_L(x^{q_t}) &= \sum_{t=1}^T (f_t)_L(x^t) \\ \sum_{t=1}^{\bar{T}} \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) &= \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \end{aligned}$$

Therefore, we conclude that

$$\sum_{t=1}^T (f_t)_L(x^t) \leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) + \frac{n}{2\eta_T} + \frac{L^2}{2} \left( \sum_{t=1}^T \eta_t \right) + L^2 \left( \sum_{t=1}^T \sum_{s=q_t}^{t-1} \eta_s \right)$$

Taking the expectation of both sides yields the desired inequality.  $\square$

*Proof of Lemma 7.* Using the same argument as in Lemma 2, we have

$$(x^t - x)^\top \hat{g}^{q_t} \leq \frac{1}{2\eta_t} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta_t}{2} \|\hat{g}^{q_t}\|^2$$

Since our pooling policy is induced by a priority queue,  $\hat{g}^{q_t}$  has never been used before updating  $x^{t+1}$ . Thus, we have  $\mathbb{E}[\hat{g}^{q_t} \mid x^t] = \mathbb{E}[\hat{g}^{q_t} \mid x^{q_t}]$  and  $\mathbb{E}[\|\hat{g}^{q_t}\|^2 \mid x^t] = \mathbb{E}[\|\hat{g}^{q_t}\|^2 \mid x^{q_t}]$ . By Lemma 3, we have  $\mathbb{E}[\hat{g}^{q_t} \mid x^{q_t}] = g^{q_t}$  and  $\mathbb{E}[\|\hat{g}^{q_t}\|^2 \mid x^{q_t}] \leq \frac{8n^2 L^2}{\mu_{q_t}}$  for all  $t \geq 1$ . Putting these pieces together yields that

$$(x^t - x)^\top g^{q_t} \leq \frac{1}{2\eta_t} (\|x - x^t\|^2 - \mathbb{E}[\|x - x^{t+1}\|^2 \mid x^t]) + \frac{4n^2 L^2 \eta_t}{\mu_{q_t}}$$

Since  $f_t = \bar{f}_t - \underline{f}_t$  where  $\bar{f}_t$  and  $\underline{f}_t$  are both normalized and non-decreasing,  $\bar{f}_t$  is  $\alpha$ -weakly DR-submodular and  $\underline{f}_t$  is  $\beta$ -weakly DR-supermodular, Proposition 1 implies that

$$(x^{q_t} - x)^\top g^{q_t} \geq (f_{q_t})_L(x^{q_t}) - \left( \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) \right)$$

By Lemma 1, we have  $\|g^t\| \leq L$  for all  $t \geq 1$ . Then, we have

$$(f_{q_t})_L(x^{q_t}) \leq \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) + \frac{1}{2\eta_t} (\|x - x^t\|^2 - \mathbb{E}[\|x - x^{t+1}\|^2 \mid x^t]) + L\|x^{q_t} - x^t\| + \frac{4n^2 L^2 \eta_t}{\mu_{q_t}} \quad (26)$$

Further, by Lemma 3, we have

$$\|x^{q_t} - x^t\| \leq \sum_{s=q_t}^{t-1} \eta_s \|\hat{g}^s\| \leq 2(n+1)L \left( \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \quad (27)$$

Plugging Eq. (27) into Eq. (26) yields that

$$\begin{aligned} (f_{q_t})_L(x^{q_t}) &\leq \frac{1}{\alpha} (\bar{f}_{q_t})_C(x) + \beta (-\underline{f}_{q_t})_C(x) \\ &\quad + \frac{1}{2\eta_t} (\|x - x^t\|^2 - \mathbb{E}[\|x - x^{t+1}\|^2 \mid x^t]) + \frac{4n^2 L^2 \eta_t}{\mu_{q_t}} + 4nL^2 \left( \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \end{aligned}$$

By using the same argument as in Lemma 6, we have

$$\begin{aligned} \sum_{t=1}^T (f_t)_L(x^t) &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} (\bar{f}_t)_C(x) + \beta (-\underline{f}_t)_C(x) \right) + \sum_{t=1}^{\bar{T}} \frac{1}{2\eta_t} (\|x - x^t\|^2 - \mathbb{E}[\|x - x^{t+1}\|^2 \mid x^t]) \\ &\quad + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \end{aligned}$$

Taking the expectation of both sides of the above inequality and using  $\|x^t - x\| \leq \sqrt{n}$  for all  $t \geq 1$  (cf. Lemma 3) and that  $\{\eta_t\}_{t \geq 1}$  is nonincreasing yields the desired inequality.  $\square$

*Proof of Lemma 8.* Using the same argument as in Lemma 7, we have

$$(x^t - x)^\top \hat{g}^{q_t} \leq \frac{1}{2\eta_t} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta_t}{2} \|\hat{g}^{q_t}\|^2$$



and

$$(x^{q_t} - x)^\top g^{q_t} \geq (f_{q_t})_L(x^{q_t}) - \left( \frac{1}{\alpha}(\bar{f}_{q_t})_C(x) + \beta(-\underline{f}_{q_t})_C(x) \right)$$

For simplicity, we define  $e_t = \hat{g}^t - g^t$ . By Lemma 1, we have  $\|g^t\| \leq L$  for all  $t \geq 1$ . Then, we have

$$\begin{aligned} & (f_{q_t})_L(x^{q_t}) - \left( \frac{1}{\alpha}(\bar{f}_{q_t})_C(x) + \beta(-\underline{f}_{q_t})_C(x) \right) \\ & \leq (x - x^t)^\top e_{q_t} + \frac{1}{2\eta_t} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + L\|x^{q_t} - x^t\| + \frac{\eta_t}{2} \|\hat{g}^{q_t}\|^2 \end{aligned} \quad (28)$$

Plugging Eq. (27) into Eq. (28) yields that

$$\begin{aligned} & (f_{q_t})_L(x^{q_t}) - \left( \frac{1}{\alpha}(\bar{f}_{q_t})_C(x) + \beta(-\underline{f}_{q_t})_C(x) \right) \\ & \leq (x - x^t)^\top e_{q_t} + \frac{1}{2\eta_t} (\|x - x^t\|^2 - \|x - x^{t+1}\|^2) + \frac{\eta_t}{2} (\|\hat{g}^t\|^2 - \mathbb{E}[\|\hat{g}^t\|^2 \mid x^t]) + \frac{4n^2 L^2 \eta_t}{\mu_{q_t}} + 4nL^2 \left( \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \end{aligned}$$

By using the same argument as in Lemma 6, we have

$$\begin{aligned} \sum_{t=1}^T (f_t)_L(x^t) & \leq \left( \sum_{t=1}^T \frac{1}{\alpha}(\bar{f}_t)_C(x) + \beta(-\underline{f}_t)_C(x) \right) + \sum_{t=1}^{\bar{T}} (x - x^t)^\top e_{q_t} + \sum_{t=1}^{\bar{T}} \frac{\eta_t}{2} (\|\hat{g}^t\|^2 - \mathbb{E}[\|\hat{g}^t\|^2 \mid x^t]) \\ & \quad + \frac{n}{2\eta_{\bar{T}}} + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \end{aligned}$$

By the definition of the convex closure, we obtain that the convex closure of a set function  $f$  agrees with  $f$  on all the integer points [Dughmi(2009), Page 4, Proposition 3.3]. Letting  $S \subseteq [n]$ , we have  $(\bar{f}_t)_C(\chi_S) = f_t(S)$  and  $(-\underline{f}_t)_C(\chi_S) = -\beta f_t(S)$  which implies that

$$\frac{1}{\alpha}(\bar{f}_t)_C(\chi_S) + \beta(-\underline{f}_t)_C(\chi_S) = \frac{1}{\alpha} f_t(S) - \beta f_t(S)$$

Letting  $x = \chi_S$ , we have

$$\begin{aligned} \sum_{t=1}^T (f_t)_L(x^t) & \leq \underbrace{\left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right)}_{\text{I}} + \underbrace{\sum_{t=1}^{\bar{T}} (\chi_S - x^t)^\top e_{q_t} + \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{2} (\|\hat{g}^t\|^2 - \mathbb{E}[\|\hat{g}^t\|^2 \mid x^t]) \right)}_{\text{II}} \\ & \quad + \frac{n}{2\eta_{\bar{T}}} + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \end{aligned} \quad (29)$$

In what follows, we prove the high probability bounds for the terms **I** and **II** in the above inequality.

**Bounding I.** Consider the random variables  $X_t = (x^t)^\top \hat{g}^{q_t}$  for all  $1 \leq t \leq \bar{T}$  that are adapted to the natural filtration generated by the iterates  $\{x_t\}_{t \geq 1}$ . By Lemma 3 and the Hölder's inequality, we have

$$|X_t| \leq \|\hat{g}^{q_t}\|_1 \|x^t\|_\infty \leq \frac{2(n+1)L}{\mu_t}$$

Since  $\mu = \frac{n}{t^{(1-\gamma)/3}}$ , we have  $|X_t| \leq 4L\bar{T}^{\frac{1-\gamma}{3}}$  for all  $1 \leq t \leq \bar{T}$ . Further, we have

$$\mathbb{E}[X_t^2 \mid x_t] \leq \mathbb{E}[\|\hat{g}^t\|_1^2 \|x^t\|_\infty^2 \mid x_t] \leq \frac{2(n+1)^2 L^2}{\mu_t} \leq 8nL^2 \bar{T}^{\frac{1-\gamma}{3}}$$

Since  $\mathbb{E}[\hat{g}^{q_t} \mid x^t] = g^{q_t}$  and  $e_{q_t} = \hat{g}^{q_t} - g^{q_t}$ , Proposition 3 implies that

$$\mathbb{P} \left( \left| \sum_{t=1}^{\bar{T}} (x^t)^\top e_{q_t} \right| > 4L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n \log \frac{1}{\delta}} + 2L\bar{T}^{\frac{1-\gamma}{3}} \log \frac{1}{\delta} \right) \leq \delta$$

Since  $\bar{T} \geq T > \log^{\frac{3}{2+\gamma}}(1/\delta)$ , we have  $\bar{T}^{\frac{4-\gamma}{6}} \sqrt{\log \frac{1}{\delta}} \geq \bar{T}^{\frac{1-\gamma}{3}} \log \frac{1}{\delta}$ . This implies that

$$\mathbb{P} \left( \left| \sum_{t=1}^{\bar{T}} (x^t)^\top e_{q_t} \right| > 6L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n \log \frac{1}{\delta}} \right) \leq \delta$$

Similarly, we fix a set  $S \subseteq [n]$  and consider the random variable  $X_t = (\chi_S)^\top \hat{g}^t$  for all  $1 \leq t \leq \bar{T}$  that are adapted to the natural filtration generated by the iterates  $\{x_t\}_{t \geq 1}$ . By repeating the above argument with  $\frac{\delta}{2^n}$ , we have

$$\mathbb{P} \left( \left| \sum_{t=1}^{\bar{T}} (\chi_S)^\top e_{q_t} \right| > 6L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n \log(2^n/\delta)} \right) \leq \frac{\delta}{2^n}$$

By taking a union bound over the  $2^n$  choices of  $S$ , we obtain that

$$\mathbb{P} \left( \left| \sum_{t=1}^{\bar{T}} (\chi_S)^\top e_{q_t} \right| > 6L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n \log(2^n/\delta)} \right) \leq \delta, \quad \text{for any } S \subseteq [n]$$

Since  $\sqrt{n \log(2^n/\delta)} \leq \sqrt{n^2 + n \log \frac{1}{\delta}}$ , we have  $\mathbf{I} \leq 12L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n^2 + n \log \frac{1}{\delta}}$  with probability at least  $1 - 2\delta$ .

**Bounding II.** Consider the random variables  $X_t = \frac{\eta_t}{2} \|\hat{g}^{q_t}\|^2$  for all  $1 \leq t \leq \bar{T}$  that are adapted to the natural filtration generated by the iterates  $\{x^t\}_{t \geq 1}$ . By Lemma 3, we have  $|X_t| \leq \frac{(n+1)^2 L^2 \eta_t}{\mu_t^2}$ . Since  $\eta_t = \frac{1}{L t^{(2+\gamma)/3}}$  and  $\mu_t = \frac{n}{t^{(1-\gamma)/3}}$ , we have  $|X_t| \leq 4L$ . Further, we have

$$\mathbb{E}[X_t^2 \mid x_t] \leq \frac{(n+1)^4 L^4 \eta_t^2}{2\mu_t^3} \leq 8nL^2$$

Applying Proposition 3, we have

$$\mathbb{P} \left( \left| \sum_{t=1}^{\bar{T}} \frac{\eta_t}{2} (\|\hat{g}^{q_t}\|^2 - \mathbb{E}[\|\hat{g}^{q_t}\|^2 \mid x^t]) \right| > 4L\sqrt{n\bar{T} \log \frac{1}{\delta}} + 2L \log \frac{1}{\delta} \right) \leq \delta$$

Since  $\bar{T} \geq T > \log^{\frac{3}{2+\gamma}}(1/\delta)$ , we have  $\sqrt{\bar{T} \log \frac{1}{\delta}} \geq \log \frac{1}{\delta}$ . This implies that

$$\mathbb{P} \left( \left| \sum_{t=1}^{\bar{T}} \frac{\eta_t}{2} (\|\hat{g}^{q_t}\|^2 - \mathbb{E}[\|\hat{g}^{q_t}\|^2 \mid x^t]) \right| > 6L\sqrt{n\bar{T} \log \frac{1}{\delta}} \right) \leq \delta$$

Therefore, we conclude that  $\mathbf{II} \leq 6L\sqrt{n\bar{T}\log\frac{1}{\delta}}$  with probability at least  $1 - \delta$ . Putting these pieces together with Eq. (29) yields that

$$\begin{aligned} \sum_{t=1}^T (f_t)_L(x^t) &\leq \left( \sum_{t=1}^T \frac{1}{\alpha} \bar{f}_t(S) - \beta \underline{f}_t(S) \right) + \frac{n}{2\eta_{\bar{T}}} + 4n^2 L^2 \left( \sum_{t=1}^{\bar{T}} \frac{\eta_t}{\mu_{q_t}} \right) + 4nL^2 \left( \sum_{t=1}^{\bar{T}} \sum_{s=q_t}^{t-1} \frac{\eta_s}{\mu_s} \right) \\ &\quad + 12L\bar{T}^{\frac{4-\gamma}{6}} \sqrt{n^2 + n \log \frac{1}{\delta}} + 6L\sqrt{n\bar{T}\log\frac{1}{\delta}} \end{aligned}$$

with probability at least  $1 - 3\delta$ . □